

# SoC 101:

a.k.a., "**Everything you wanted to know about a computer but were afraid to ask**"

# Communicating with Peripherals

*(How to build a router...)*

Prof. Adam Teman

EnIICS Labs, Bar-Ilan University

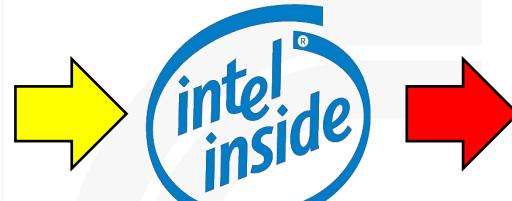
22 May 2023

Heavily based on the wonderful lecture  
“Interfaces: External/Internal, or why CPUs suck”  
by Tzachi Noy, 2019

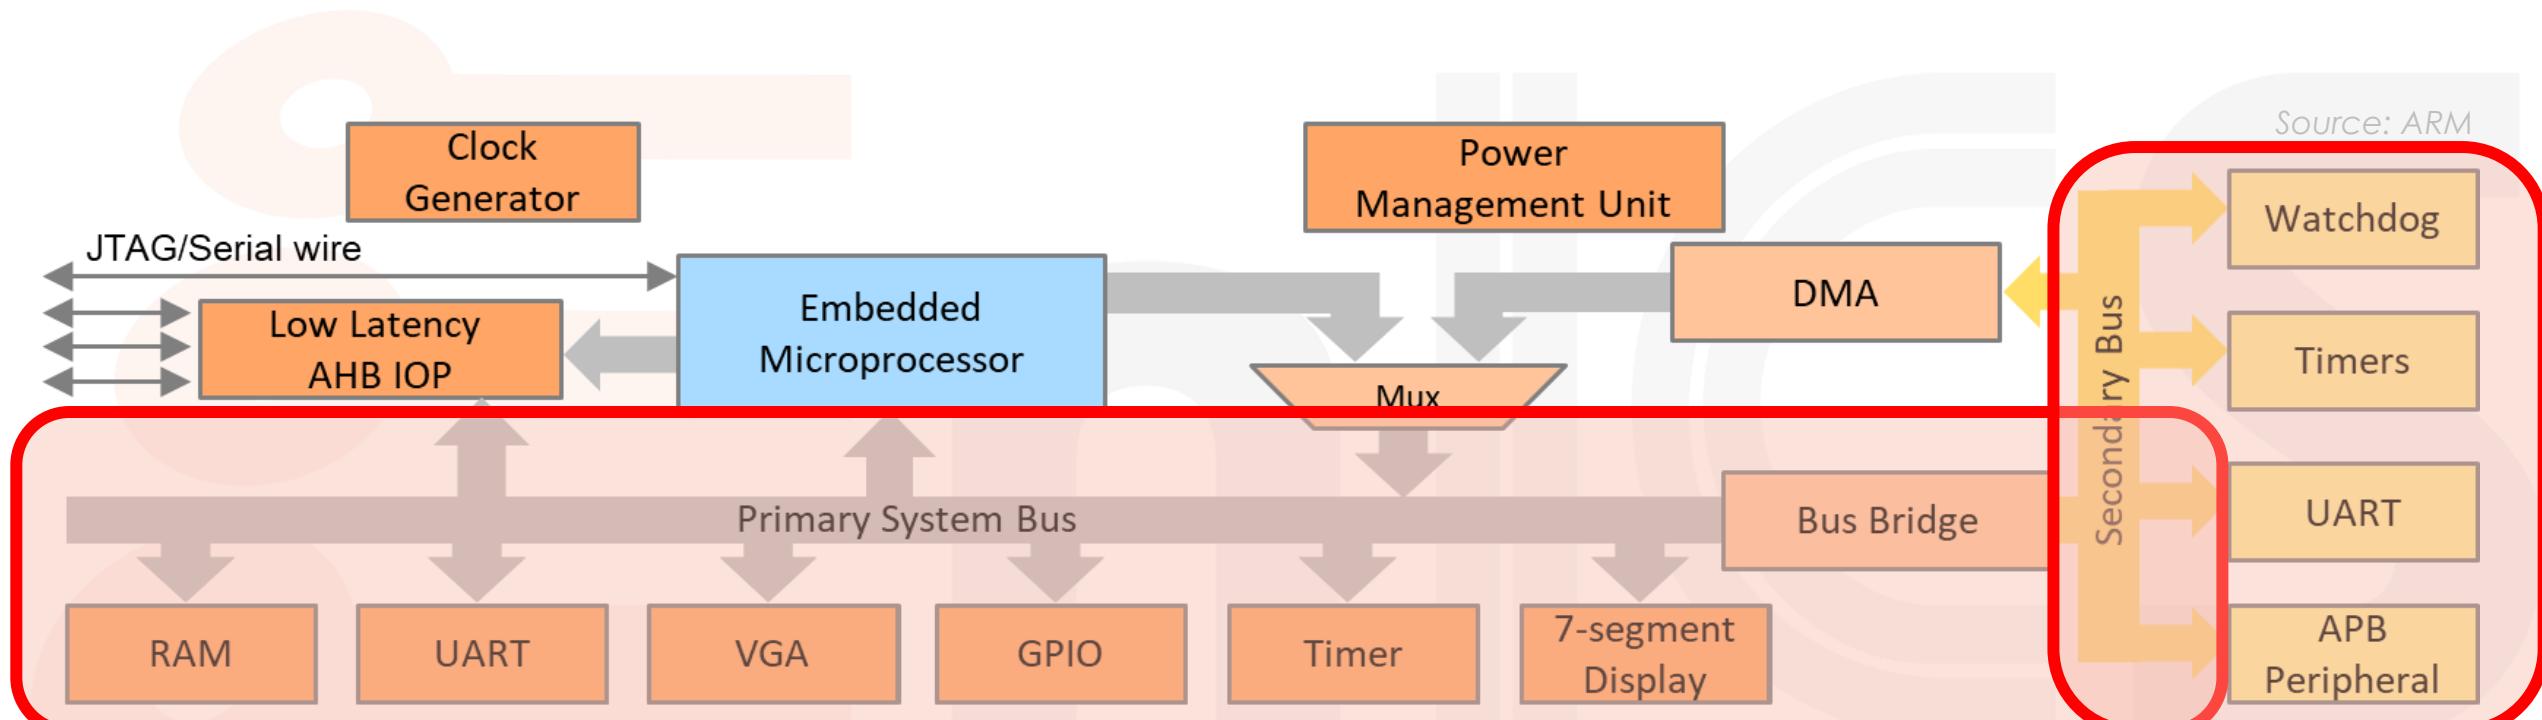
# What do a car and a router have in common?



# Both the car and the router have interfaces



# This Lecture



We will “design” a **router** as an example  
to introduce these concepts



# Lecture Outline



Communicating with the outside world



6  
Emerging Nanoscaled Integrated Circuits and Systems Labs



Offloading the CPU



13  
Emerging Nanoscaled Integrated Circuits and Systems Labs



Dealing with Faster Interfaces



20  
Emerging Nanoscaled Integrated Circuits and Systems Labs



More Offloading



31  
Emerging Nanoscaled Integrated Circuits and Systems Labs



Memory



41  
Emerging Nanoscaled Integrated Circuits and Systems Labs



Finishing our Design



48  
Emerging Nanoscaled Integrated Circuits and Systems Labs



GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

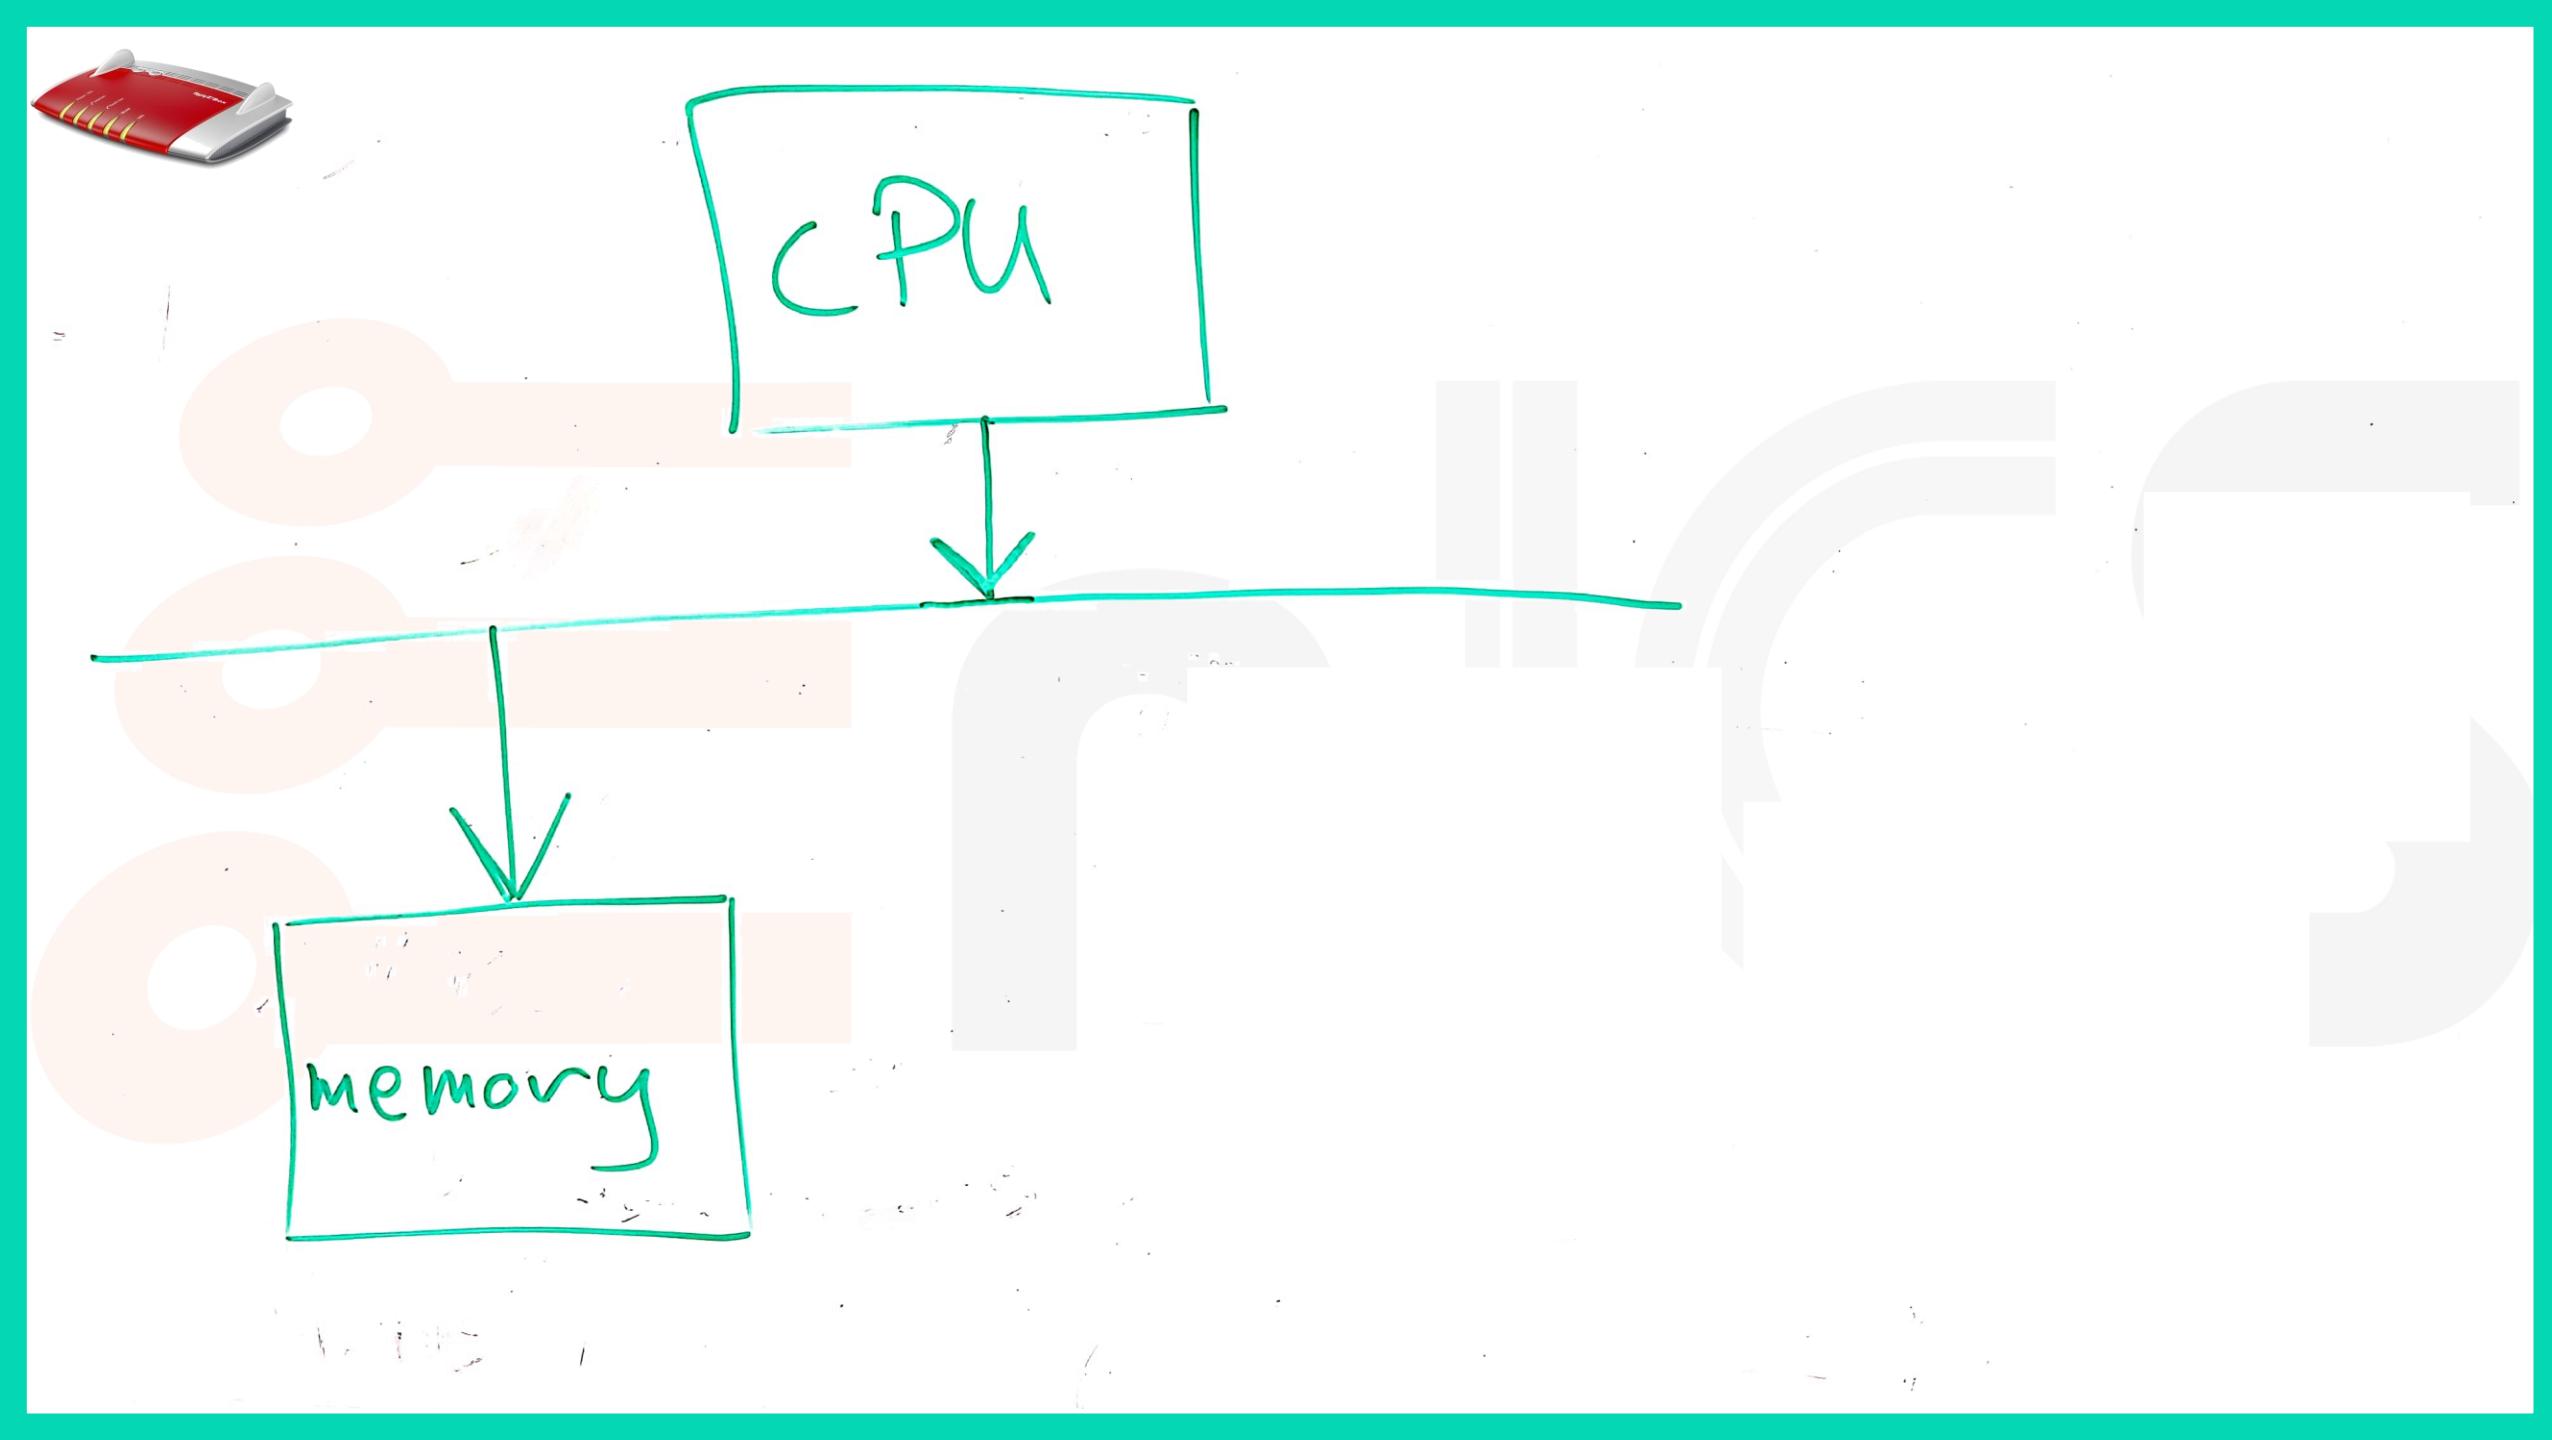
# Communicating with the outside world



Emerging Nanoscaled  
Integrated Circuits and Systems Labs

The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University





# Reminder: Memory-Mapped I/O

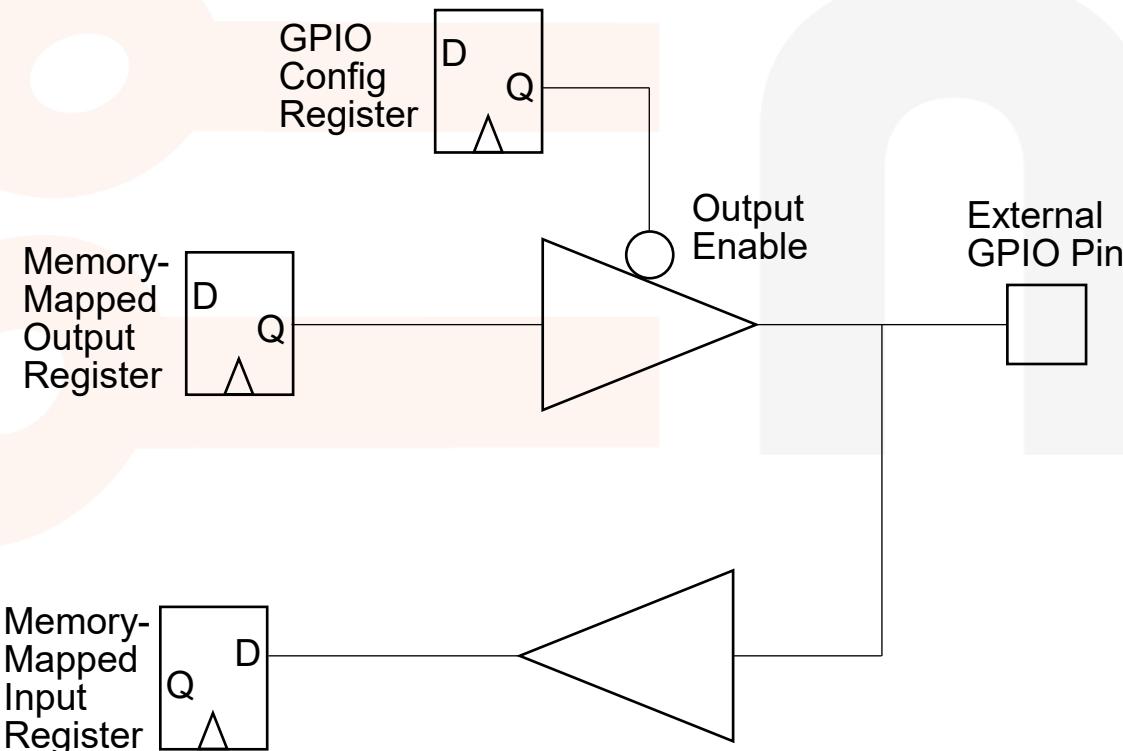
- Registers and I/O Devices are given an address in the system's memory map:
  - Everything is treated the same as memory.
  - To communicate with an I/O, we write to and read from these addresses.
  - These are achieved with simple `load` and `store` assembly commands.
- In C, we can define two functions, `peek` and `poke`, to accomplish this easily:
  - Now to access a register, just define its address, and use these functions

```
int peek (char *location) {  
    // Read from a memory-mapped address  
    return *location;  
}  
void poke (char *location, char newval) {  
    // Write to a memory-mapped address  
    (*location) = newval;  
}
```

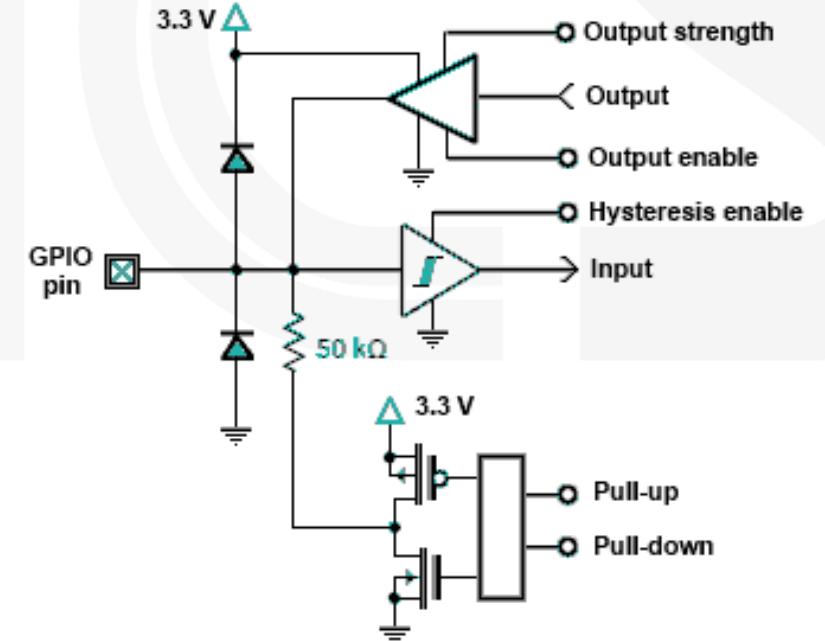
```
#define DEV1 0x1000  
...  
dev_status = peek(DEV1)  
...  
poke(DEV1,8);
```

# General Purpose I/O (GPIO)

- Most microcontrollers have a set of general purpose input/output (GPIO) pins.
  - Can be configured as input pins or output pins.
  - Can be programmed by software for various purposes.

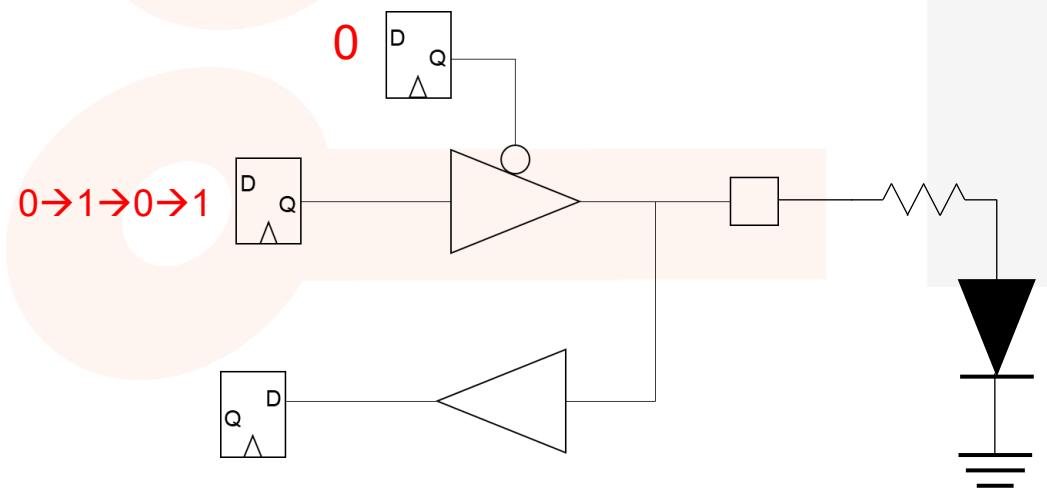


Equivalent Circuit for Raspberry Pi GPIO pins



# Example: Blinking a LED

- First, configure the **GPIO** to be an **output**.
- Next, create an **infinite loop** that:
  - Toggles the state of the output.
  - Waits for a given period.



```
#define GPIO_CONFIG_REG 0x10000000
#define GPIO_OUTPUT_REG 0x10000001
#define GPIO_BLINK_PIN 0b00000001
#define BLINK_PERIOD 1000000

int main () {
    // Set LED connected GPIO PIN to output
    toggle_config |= peek(GPIO_CONFIG_REG);
    poke(GPIO_CONFIG_REG,toggle_config);

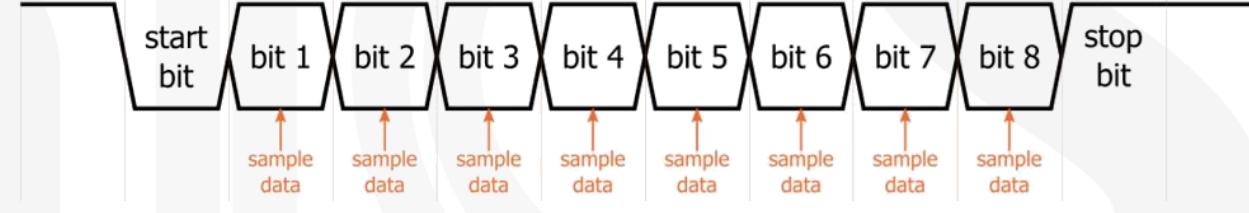
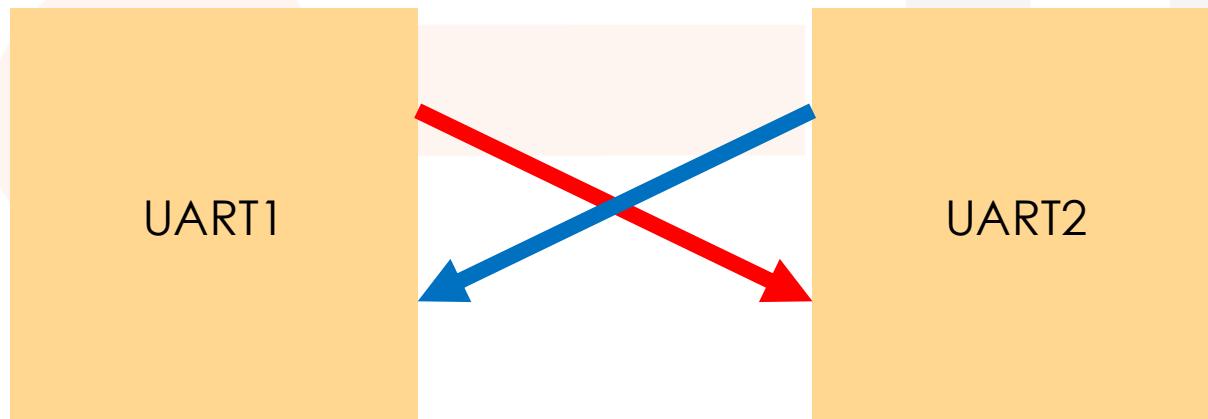
    while (true) {
        // Toggle the state of the GPIO output register
        output_status = peek(GPIO_OUTPUT_REG);
        poke(GPIO_OUTPUT_REG,
              output_status~GPIO_BLINK_PIN);
        // Wait for a predefined delay
        wait(BLINK_PERIOD);
    }
}
```

# Communicating Off-Chip

- What if we want to communicate with something *more sophisticated* than a LED or a button?
  - We need a [communication protocol](#).

- Introducing [UART](#)

- The Universal Asynchronous Receiver/Transmitter



- **Baud Rate**

- Number of bits per unit time

$$\text{Baud Rate} = \frac{1}{\text{bit time}}$$

- **Bandwidth**

$$BW = \frac{\text{data bits}}{\text{frame bits}} \cdot \text{Baud Rate}$$

- Data per unit time

# Can we use UART for our router?

**BAUD RATE:**

115,200 bits/sec

**SAMPLE RATE:**

230,400 samples/sec

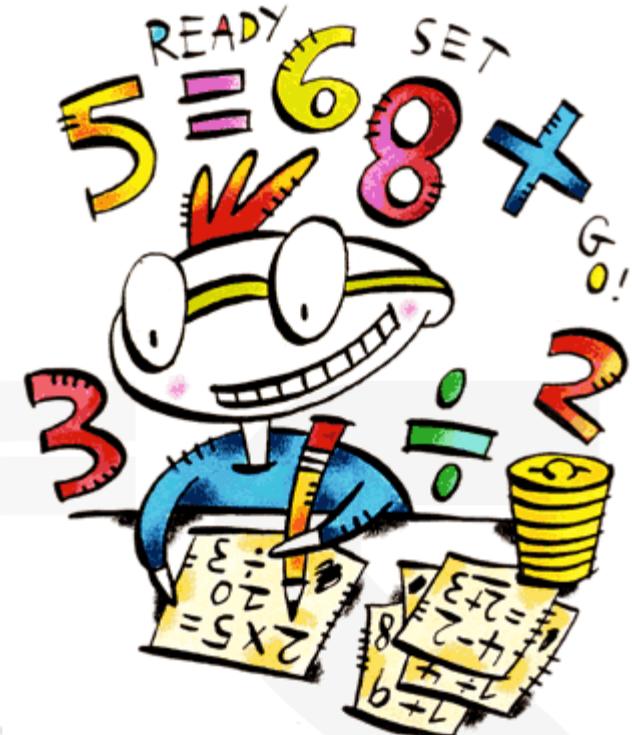
**CODE:**

40 instructions/sample

**OVERHEAD:**

9,216,000 instructions/sec

**9.2% of CPU time**



\* Assuming a 100MHz  
clock frequency

GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

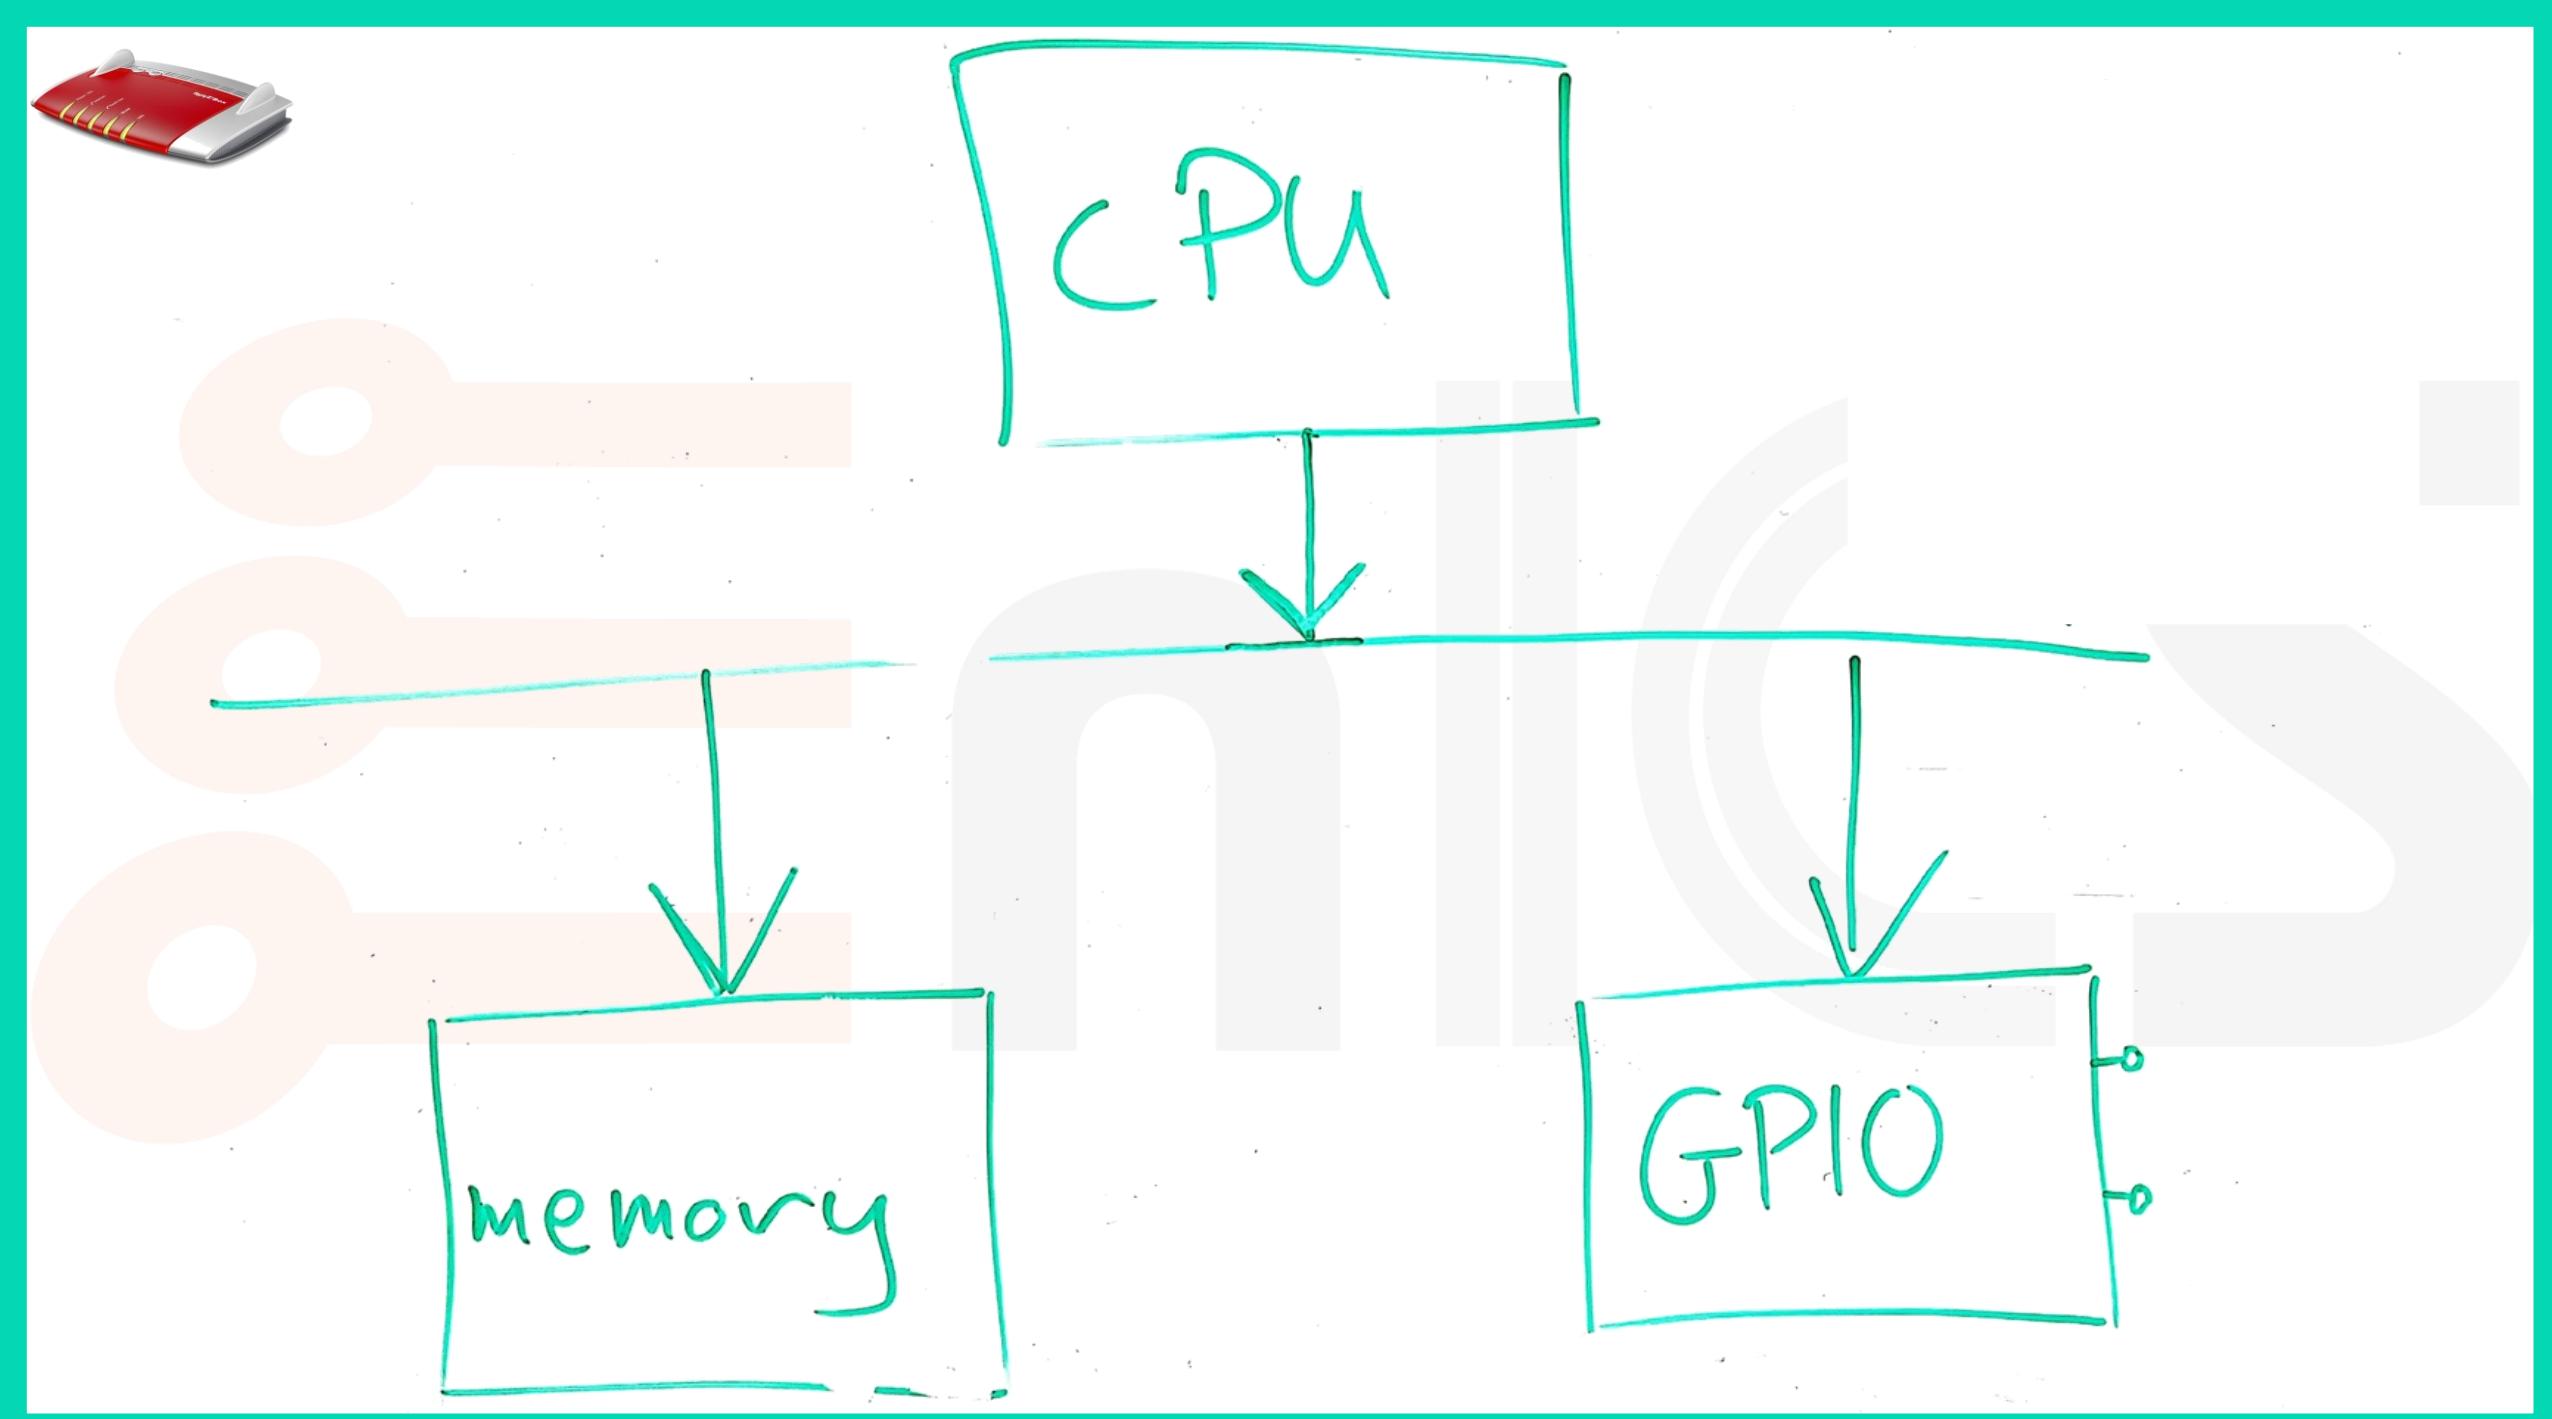
# Offloading the CPU



Emerging Nanoscaled  
Integrated Circuits and Systems Labs

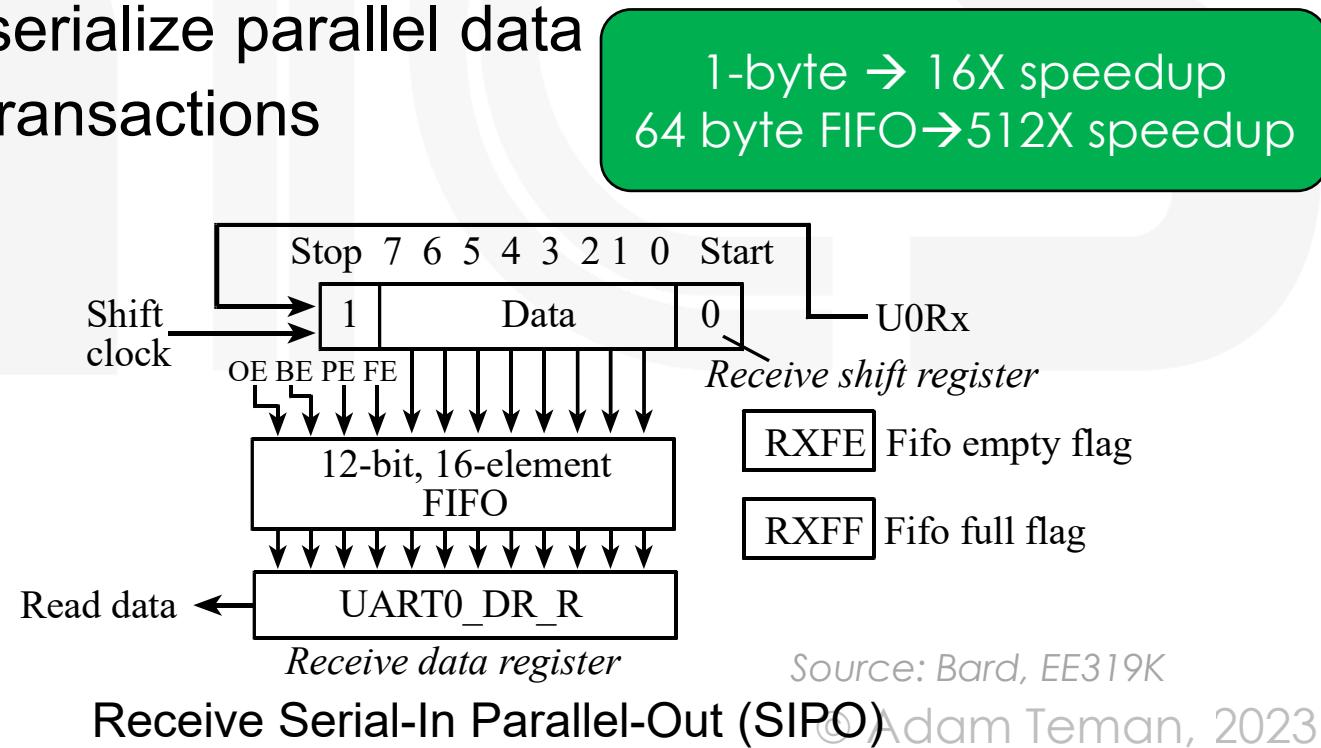
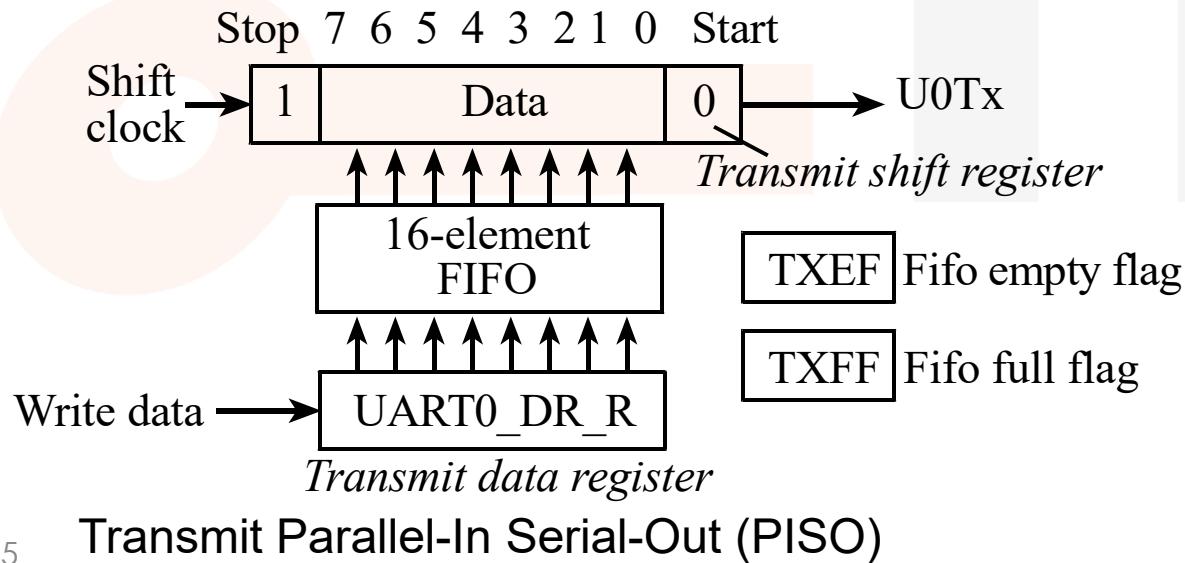
The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University





# Offload the CPU with a controller

- **UART is a slow serial protocol**
  - One bit is transferred at a time at a low **baud rate** (e.g., 1200-115200 bits/sec).
- **Integrate a specific UART controller that offloads the CPU**
  - Communicate with the **UART** through a wider register (e.g., byte, 32-bit).
  - Use a **Shift Register** to serialize/deserialize parallel data
  - Use a **FIFO** to buffer several CPU transactions



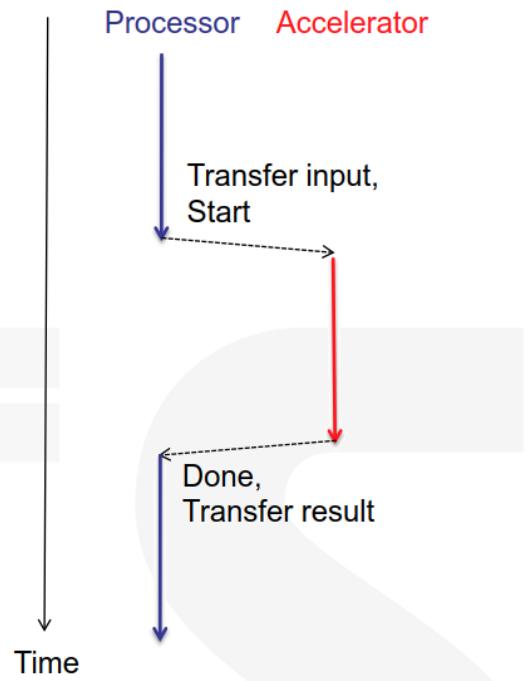
1-byte → 16X speedup  
64 byte FIFO → 512X speedup

Source: Bard, EE319K

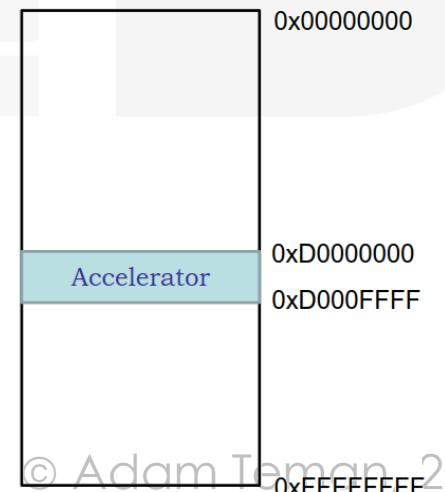
Adam Teman, 2023

# Hardware Acceleration

- CPUs are **general purpose programmable machines**
  - In other words, they are **Turing Complete**.
  - But CPUs are **not great** (sometimes **terrible**) at carrying out certain operations.
- **Offload the CPU by providing dedicated hardware**
  - The dedicated hardware can be designed to efficiently run a **specific task** → **accelerate** it.
  - The CPU can **continue running the program**, while the accelerator runs its task.
- **Data transfer achieved through memory mapping**
  - CPU writes/stores inputs and control in **memory/registers**.
  - Accelerator writes outputs/status in **memory/registers**.



$$\text{Speedup} = \frac{T_{\text{original}}}{T_{\text{unaccel}} + T_{\text{accelerated}} + T_{\text{comm}}}$$



# But how do we know when it's done?

- How can the CPU know when a new byte of data is received?

- Simple way: “**Polling**”
- Check on the status of the **UART** every so often to see if data has been received (or if it is ready to receive new data)

- **Polling** can be carried out with a “**busy-wait**” loop:

Busy-wait on input from the UART

```
while (TRUE) {  
    // Wait until a new character has been read  
    while (peek(UART_IN_STATUS)==0);  
    // Read the new character  
    achar=(char)peek(UART_DATA);  
}
```



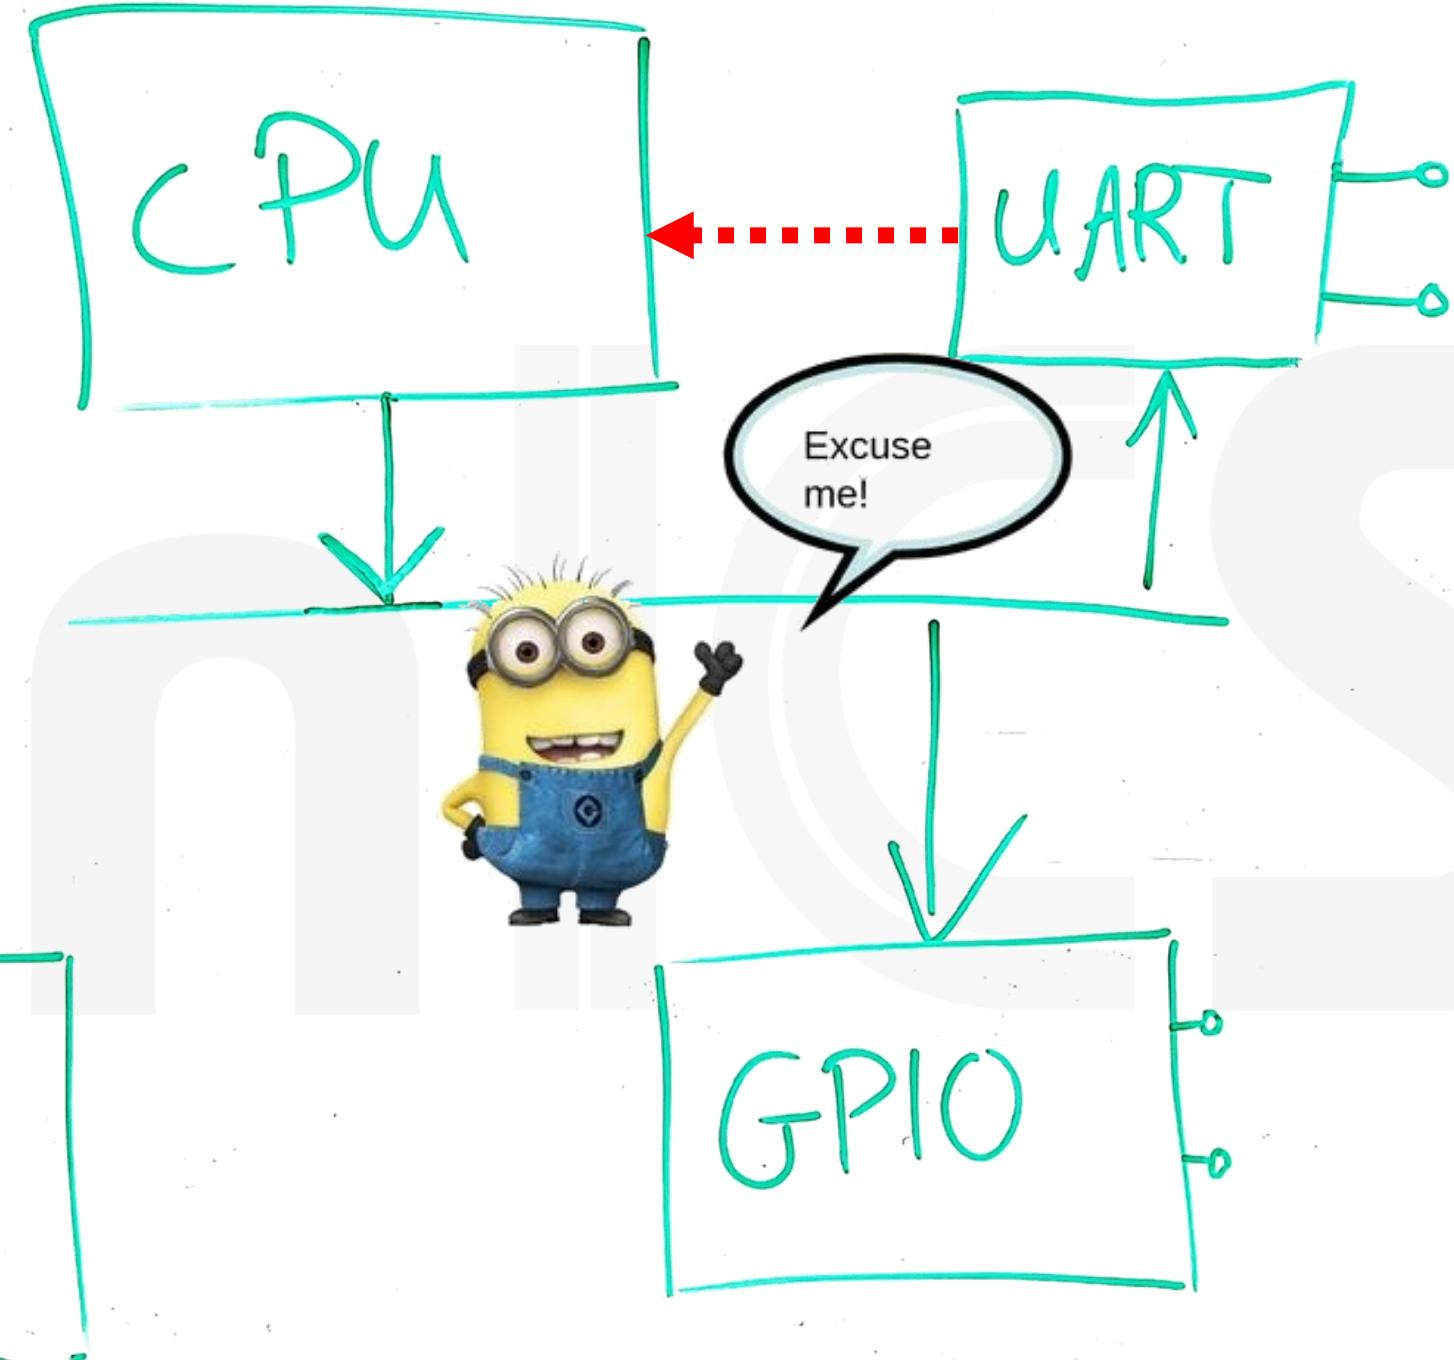
```
do {  
    // Play games  
    ...  
    // Poll to see if we're there yet.  
    status = areWeThereYet();  
} while (status == NO);
```

Busy-wait on writing to the UART

```
current_char = mystring;  
// Continue until the end of string  
while (*current_char != '\0') {  
    // Wait until the UART is ready  
    while (peek(UART_OUT_STATUS)!=0);  
    // Send character to UART  
    poke(UART_DATA_OUT, *current_char);  
    // update character pointer  
    current_char++;  
}
```



# Can't we do any better?



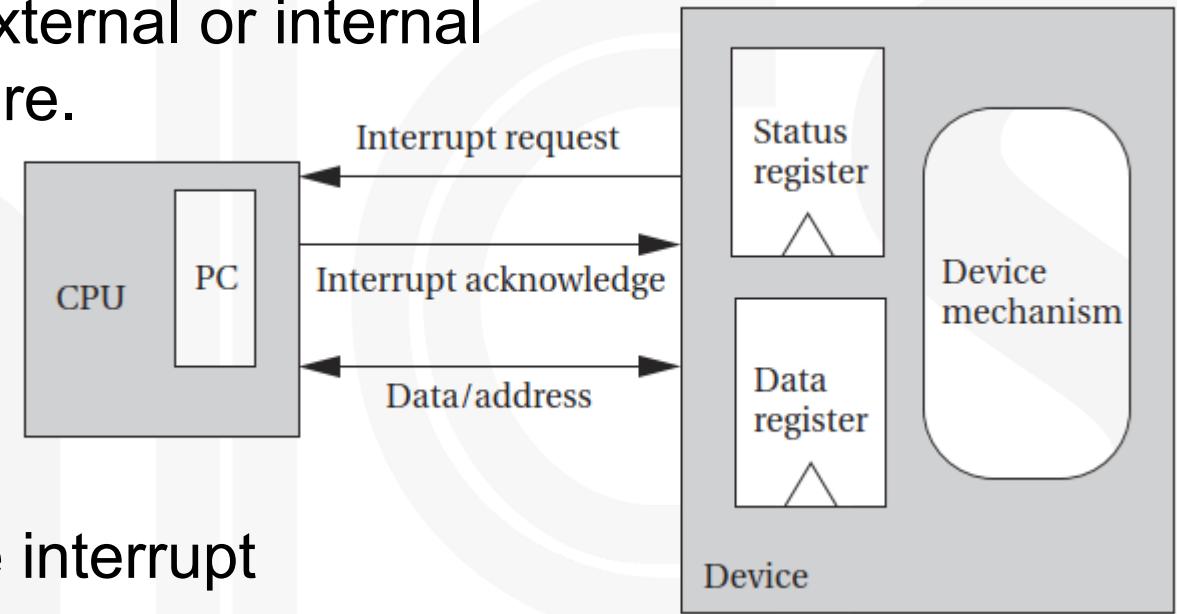
# Interrupts

- An **interrupt** is an **asynchronous signal** from a peripheral to the processor.

- Can be generated from peripherals external or internal to the processor, as well as by software.
- Frees up the CPU, while the peripheral is doing its job.

- Upon receiving an interrupt:

- The CPU decides when to *handle* the interrupt
- When ready, the CPU *acknowledges* the interrupt
- The CPU calls an **interrupt service routine (ISR)**
- Upon finishing, the ISR *returns* and the CPU continues operation



Source: Computers as Components

GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

# Dealing with Faster Interfaces



Emerging Nanoscaled  
Integrated Circuits and Systems Labs

20

The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University



# UART

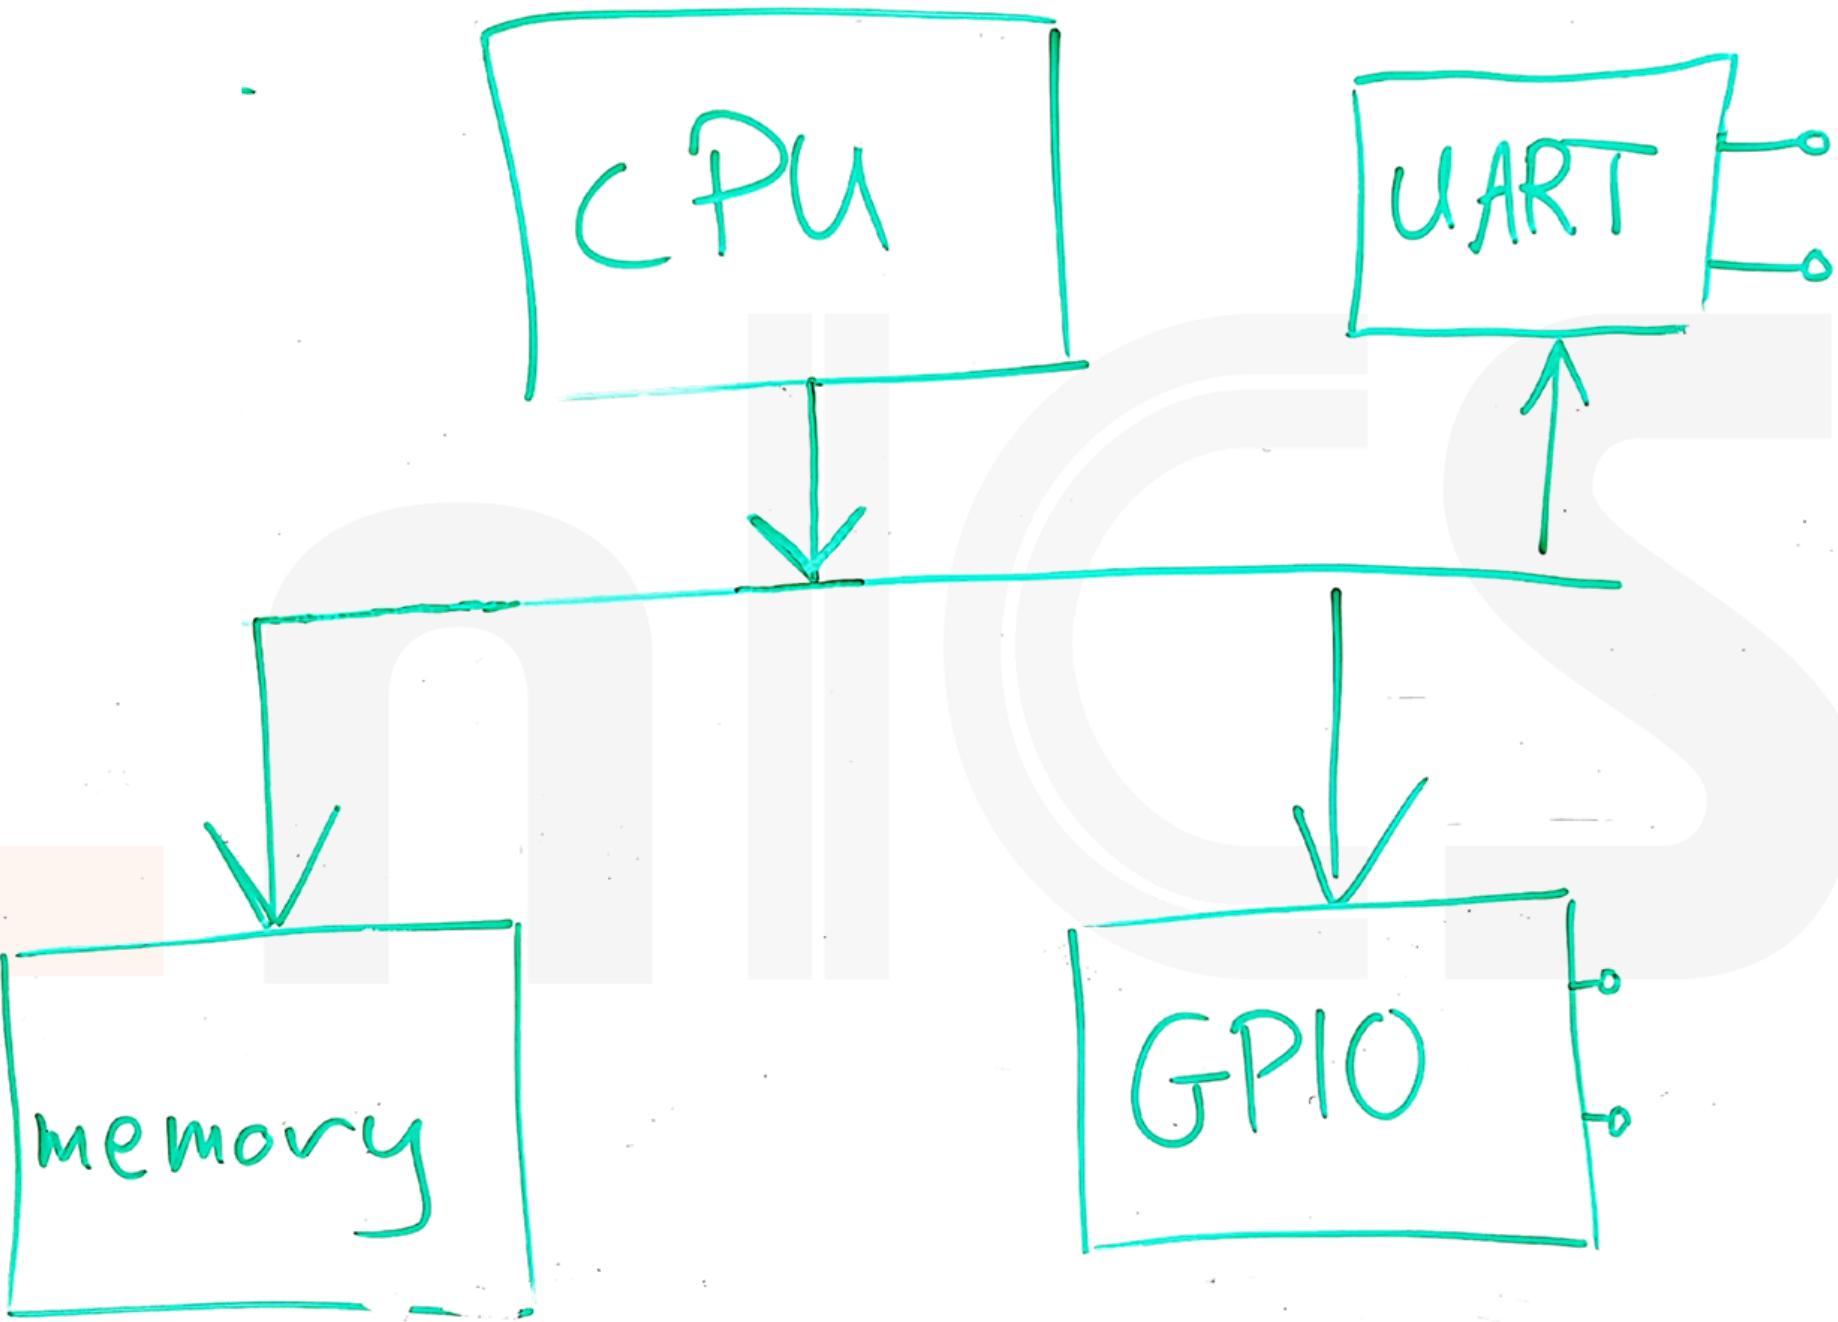
**RATE:**

115200 bit/sec  
0.1152 Mbps

**RANGE:**

15m





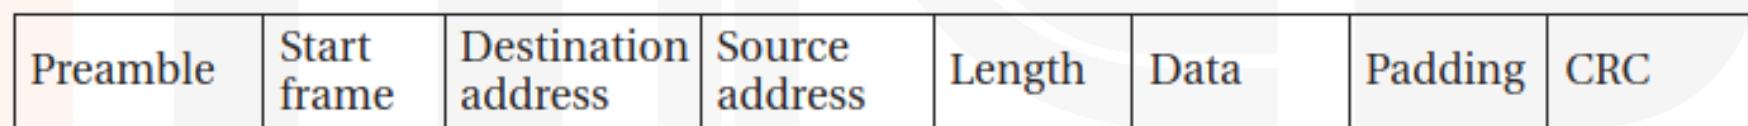
# Ethernet

- Widely used for realization of Local Area Networks (LANs)

- Bus with single signal path
- Originally: Nodes are not synchronized → *Collisions*
- Arbitration: “Carrier Sense Multiple Access with Collision Detection (CSMA/CD)”
- If collision → wait for random time → retransmit.
- Now: switched (point-to-point), fully duplex

- **Ethernet packet:**

- Addresses
- Variable-length data payload: 46 – 1518 bytes

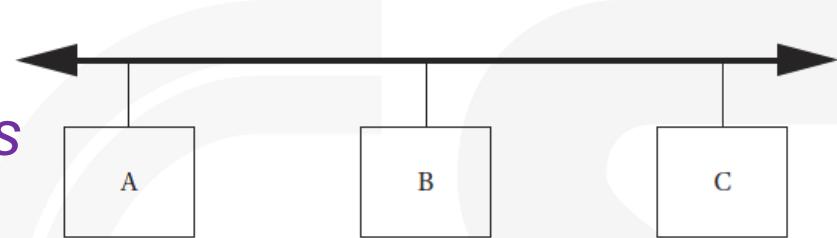


- **Throughput:**

- $10M = 2.5 \times 4\text{bit}$

- $100M = 25 \times 4\text{bit}$

- $1G = 125 \times 8\text{bit}$



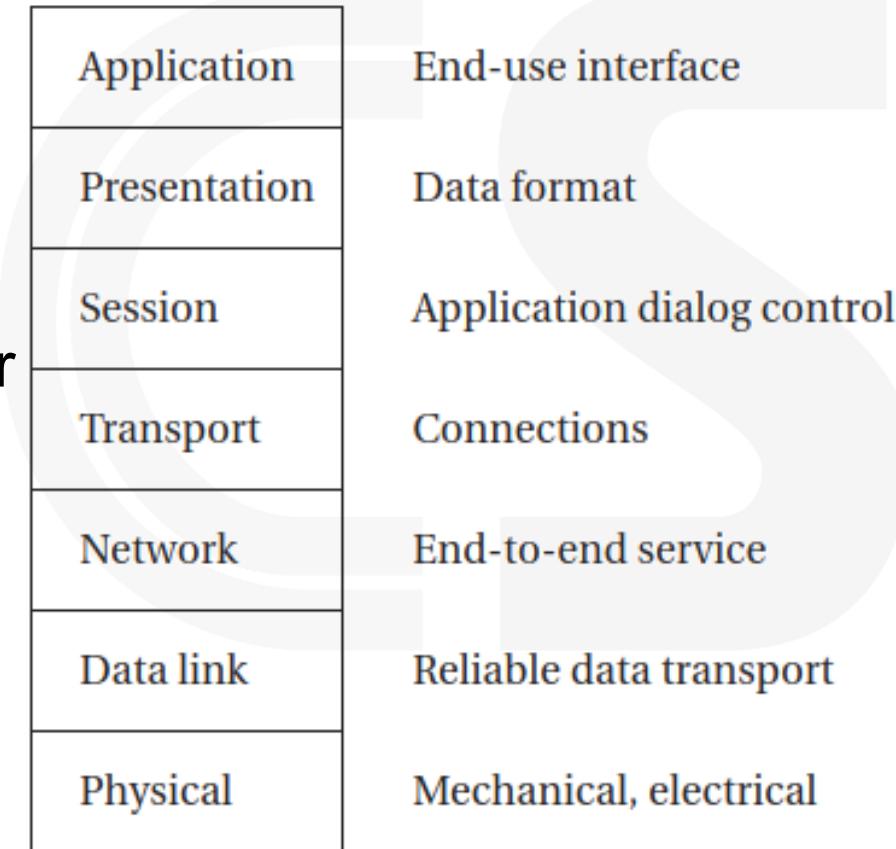
Source: Computers as Components

# Side note: The OSI Model

- The Open Systems Interconnection (OSI) model defines seven network layers.

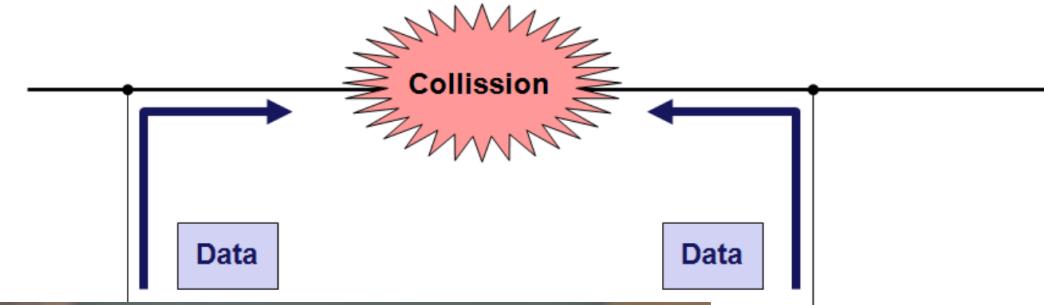
1. **Physical**: electrical and physical components
2. **Data Link**: Peer2Peer communication across a single physical layer.
3. **Network**: basic routing over the link.
4. **Transport**: ensure data is delivered in the proper order and without errors across multiple links.
5. **Session**: interaction of end-user services across a network
6. **Presentation**: defines data exchange formats including encryption and compression.
7. **Application**: interface between the network and end-user

Ethernet



Source: Computers as Components

# Ethernet



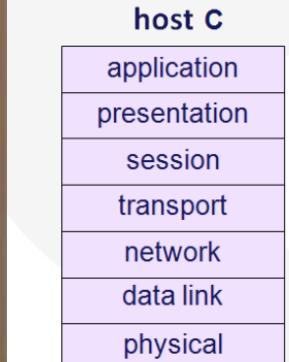
# ETHERNET PORTS



Firewall



Desktop Computer



data field: min. 46 bytes, max 1500 bytes  
© Adam Nemeth, 2023

# Let's try a simple interface: APB

- 32-bit bus
- Two phase access:
  - Address (*Setup*) phase
  - Read/Write (*Access*) phase

## 1 Setup Phase

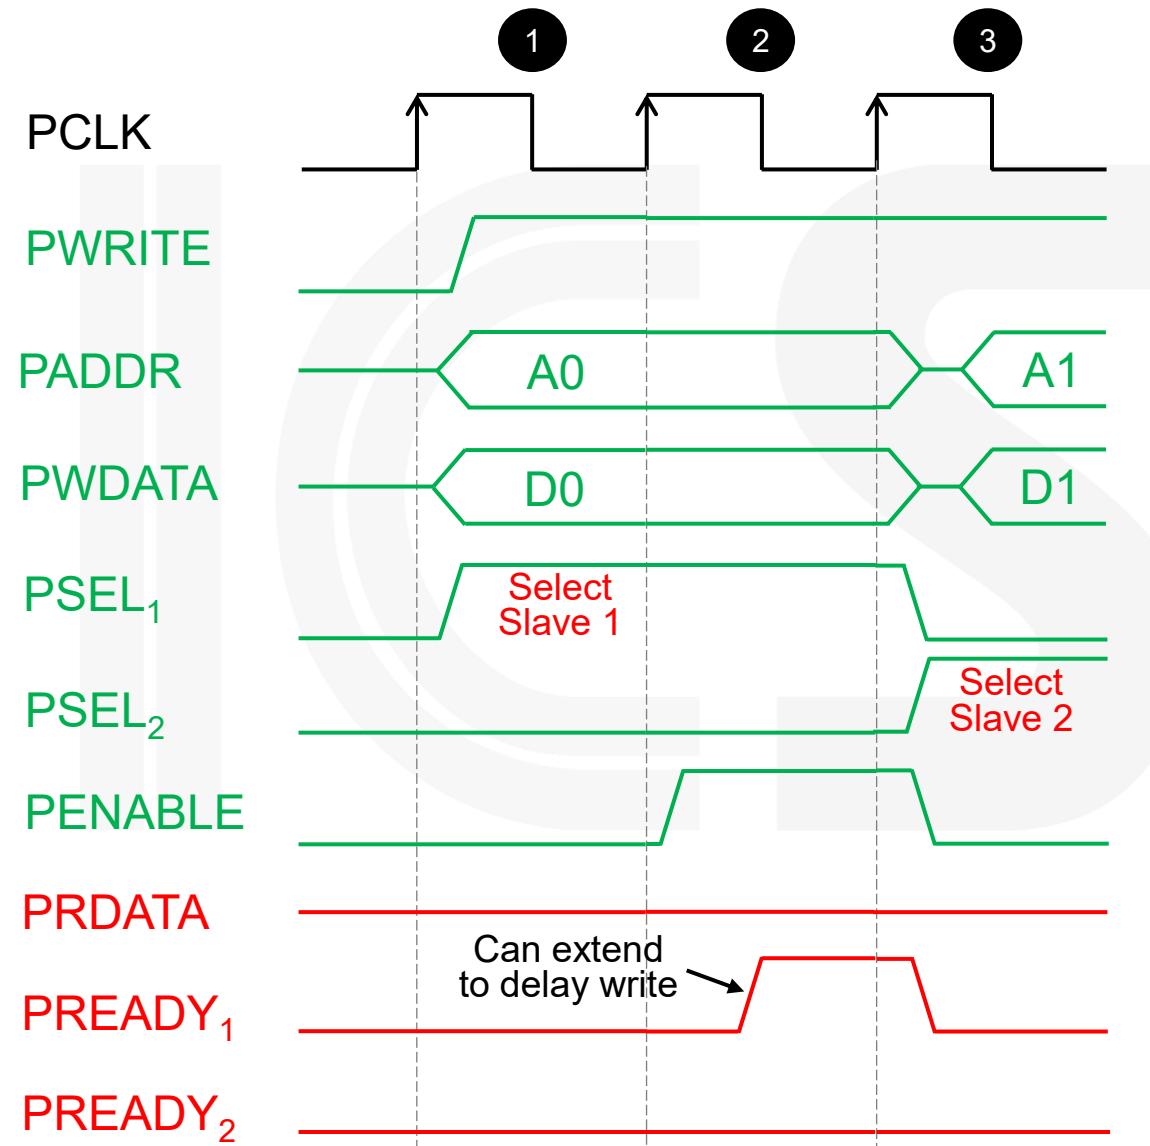
- PWRITE, PADDR, PWDATA are set
- PSEL is raised for selected Slave

## 2 Access Phase

- PENABLE is raised, with other signals held
- When selected Slave acks, PREADY is raised

## 3 Next Transfer

- PENABLE is lowered by Master
- PREADY may be lowered by Slave



# Is APB Sufficient?

**ETH RATE:**

$10^9$  bits/sec

**APB TRANSFER WIDTH:**

32 bits ↑

**APB RATE:**

2 cycles/transfer ↓

**CLOCK:**

$10^8$  cycles/sec

**APB THROUHPUT:**

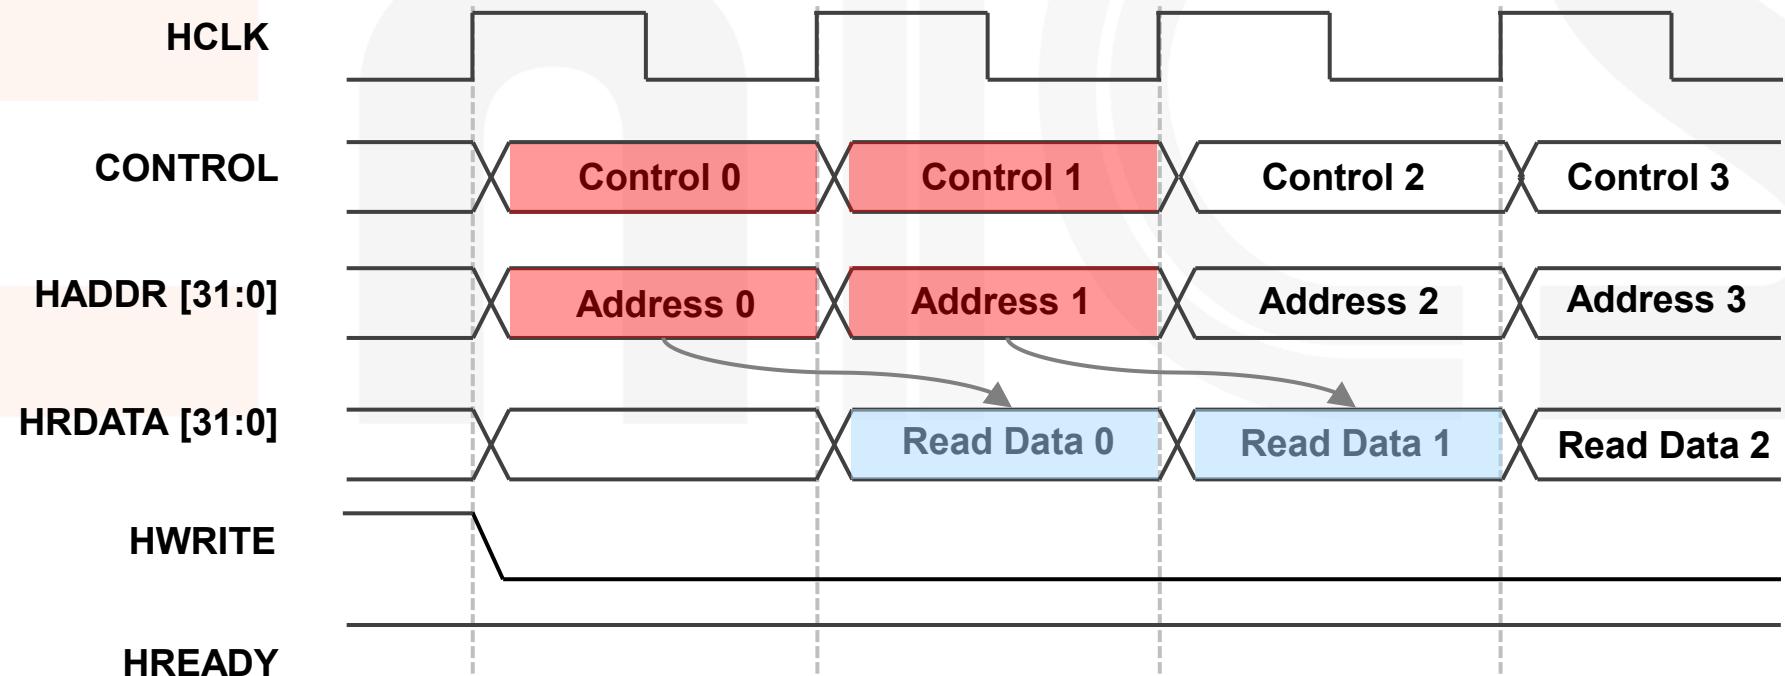
$1.6 \times 10^9$  bits/sec

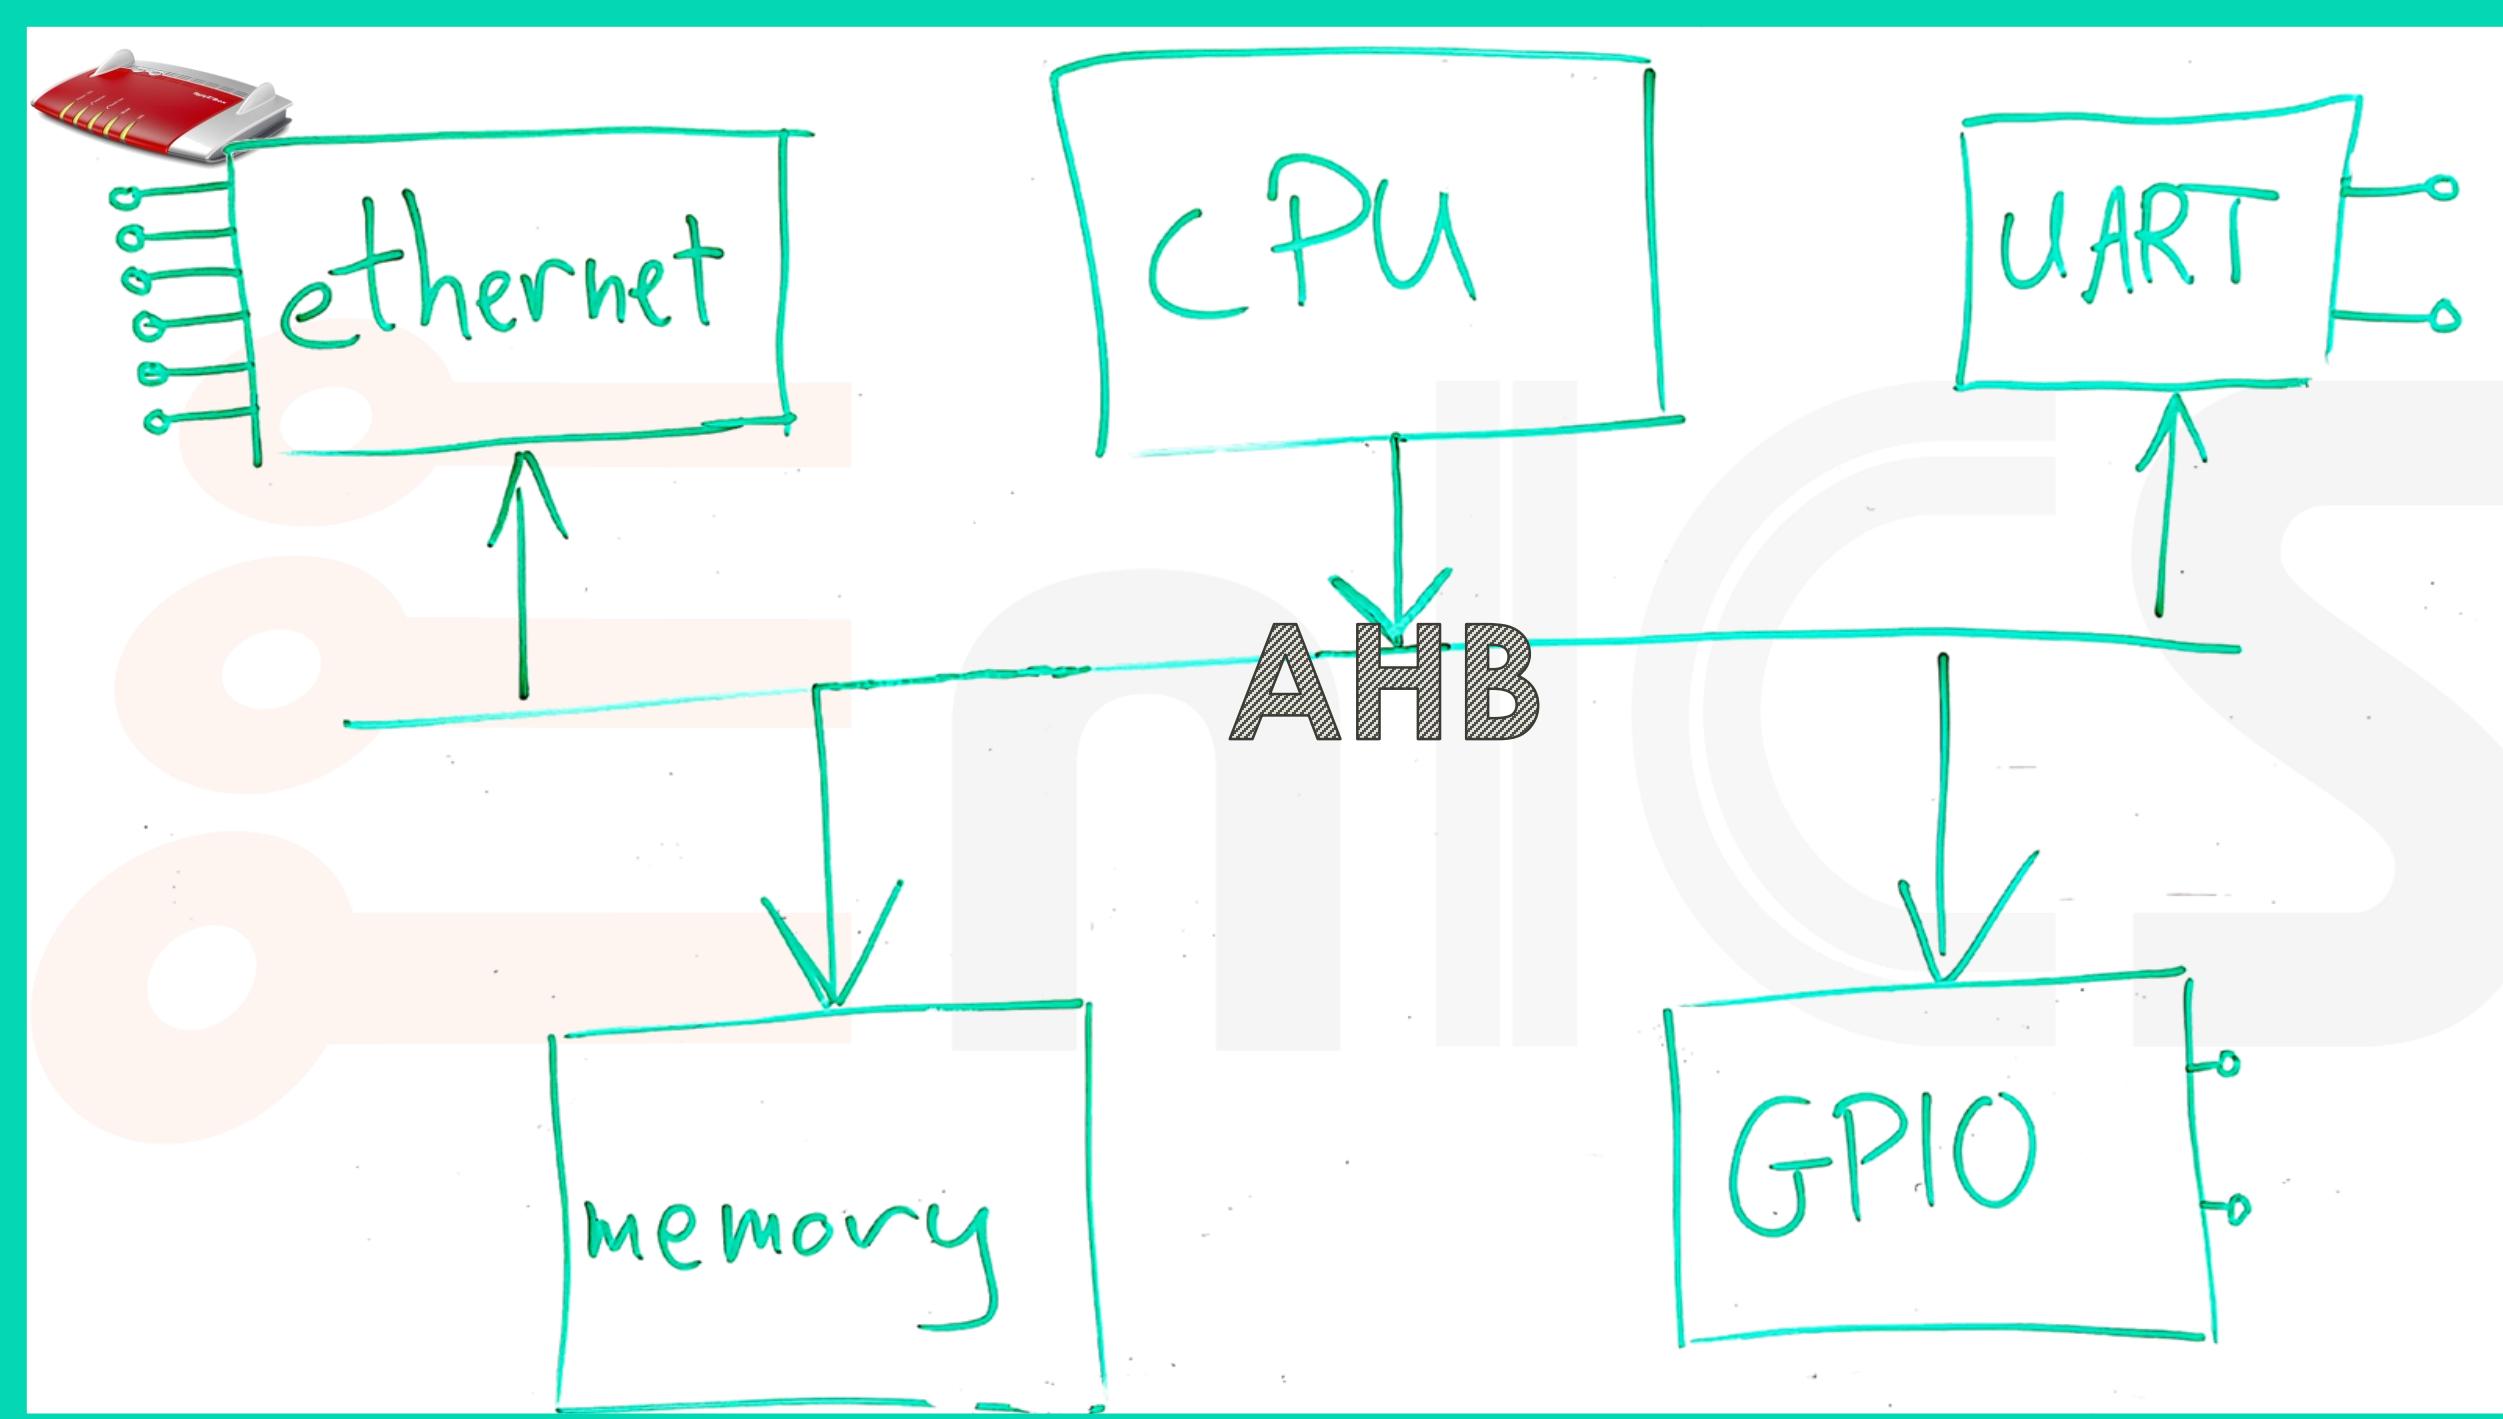
**BUT, Ethernet is FULL-DUPLEX,  
so maximum throughput is actually  
 $2 \times 10^9$  bits/sec**



# So let's make it faster: AHB

- Wider bus (>32 bits)
- Pipelined address and R/W phases (X2 throughput)
- Supports Bursts





# Is AHB fast enough?

**ETHERNET RATE:**

$2 \times 10^9$  bits/sec

**AHB TRANSFER WIDTH:** 64 bits

**AHB RATE:**

1 cycle/transfer

**CLOCK:**

$10^8$  cycles/sec

**AHB THROUGHPUT:**

$6.4 \times 10^9$  bits/sec



## BUT WHAT ABOUT THE CPU?

GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

# More Offloading



Emerging Nanoscaled  
Integrated Circuits and Systems Labs

31

The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University



# Is AHB fast enough?

**ETHERNET RATE:**

$2 \times 10^9$  bits/sec

**AHB TRANSFER WIDTH:** 64 bits

**AHB RATE:**

1 cycle/transfer

**CLOCK:**

$10^8$  cycles/sec

**AHB THROUGHPUT:**

$6.4 \times 10^9$  bits/sec



## BUT WHAT ABOUT THE CPU?

# Can the CPU support this?

**ETHERNET RATE:**

$2 \times 10^9$  bits/sec

**CPU WORD:**

32 bits

**INSTRUCTIONS PER SW/LW:**

3 inst/load

**CLOCK:**

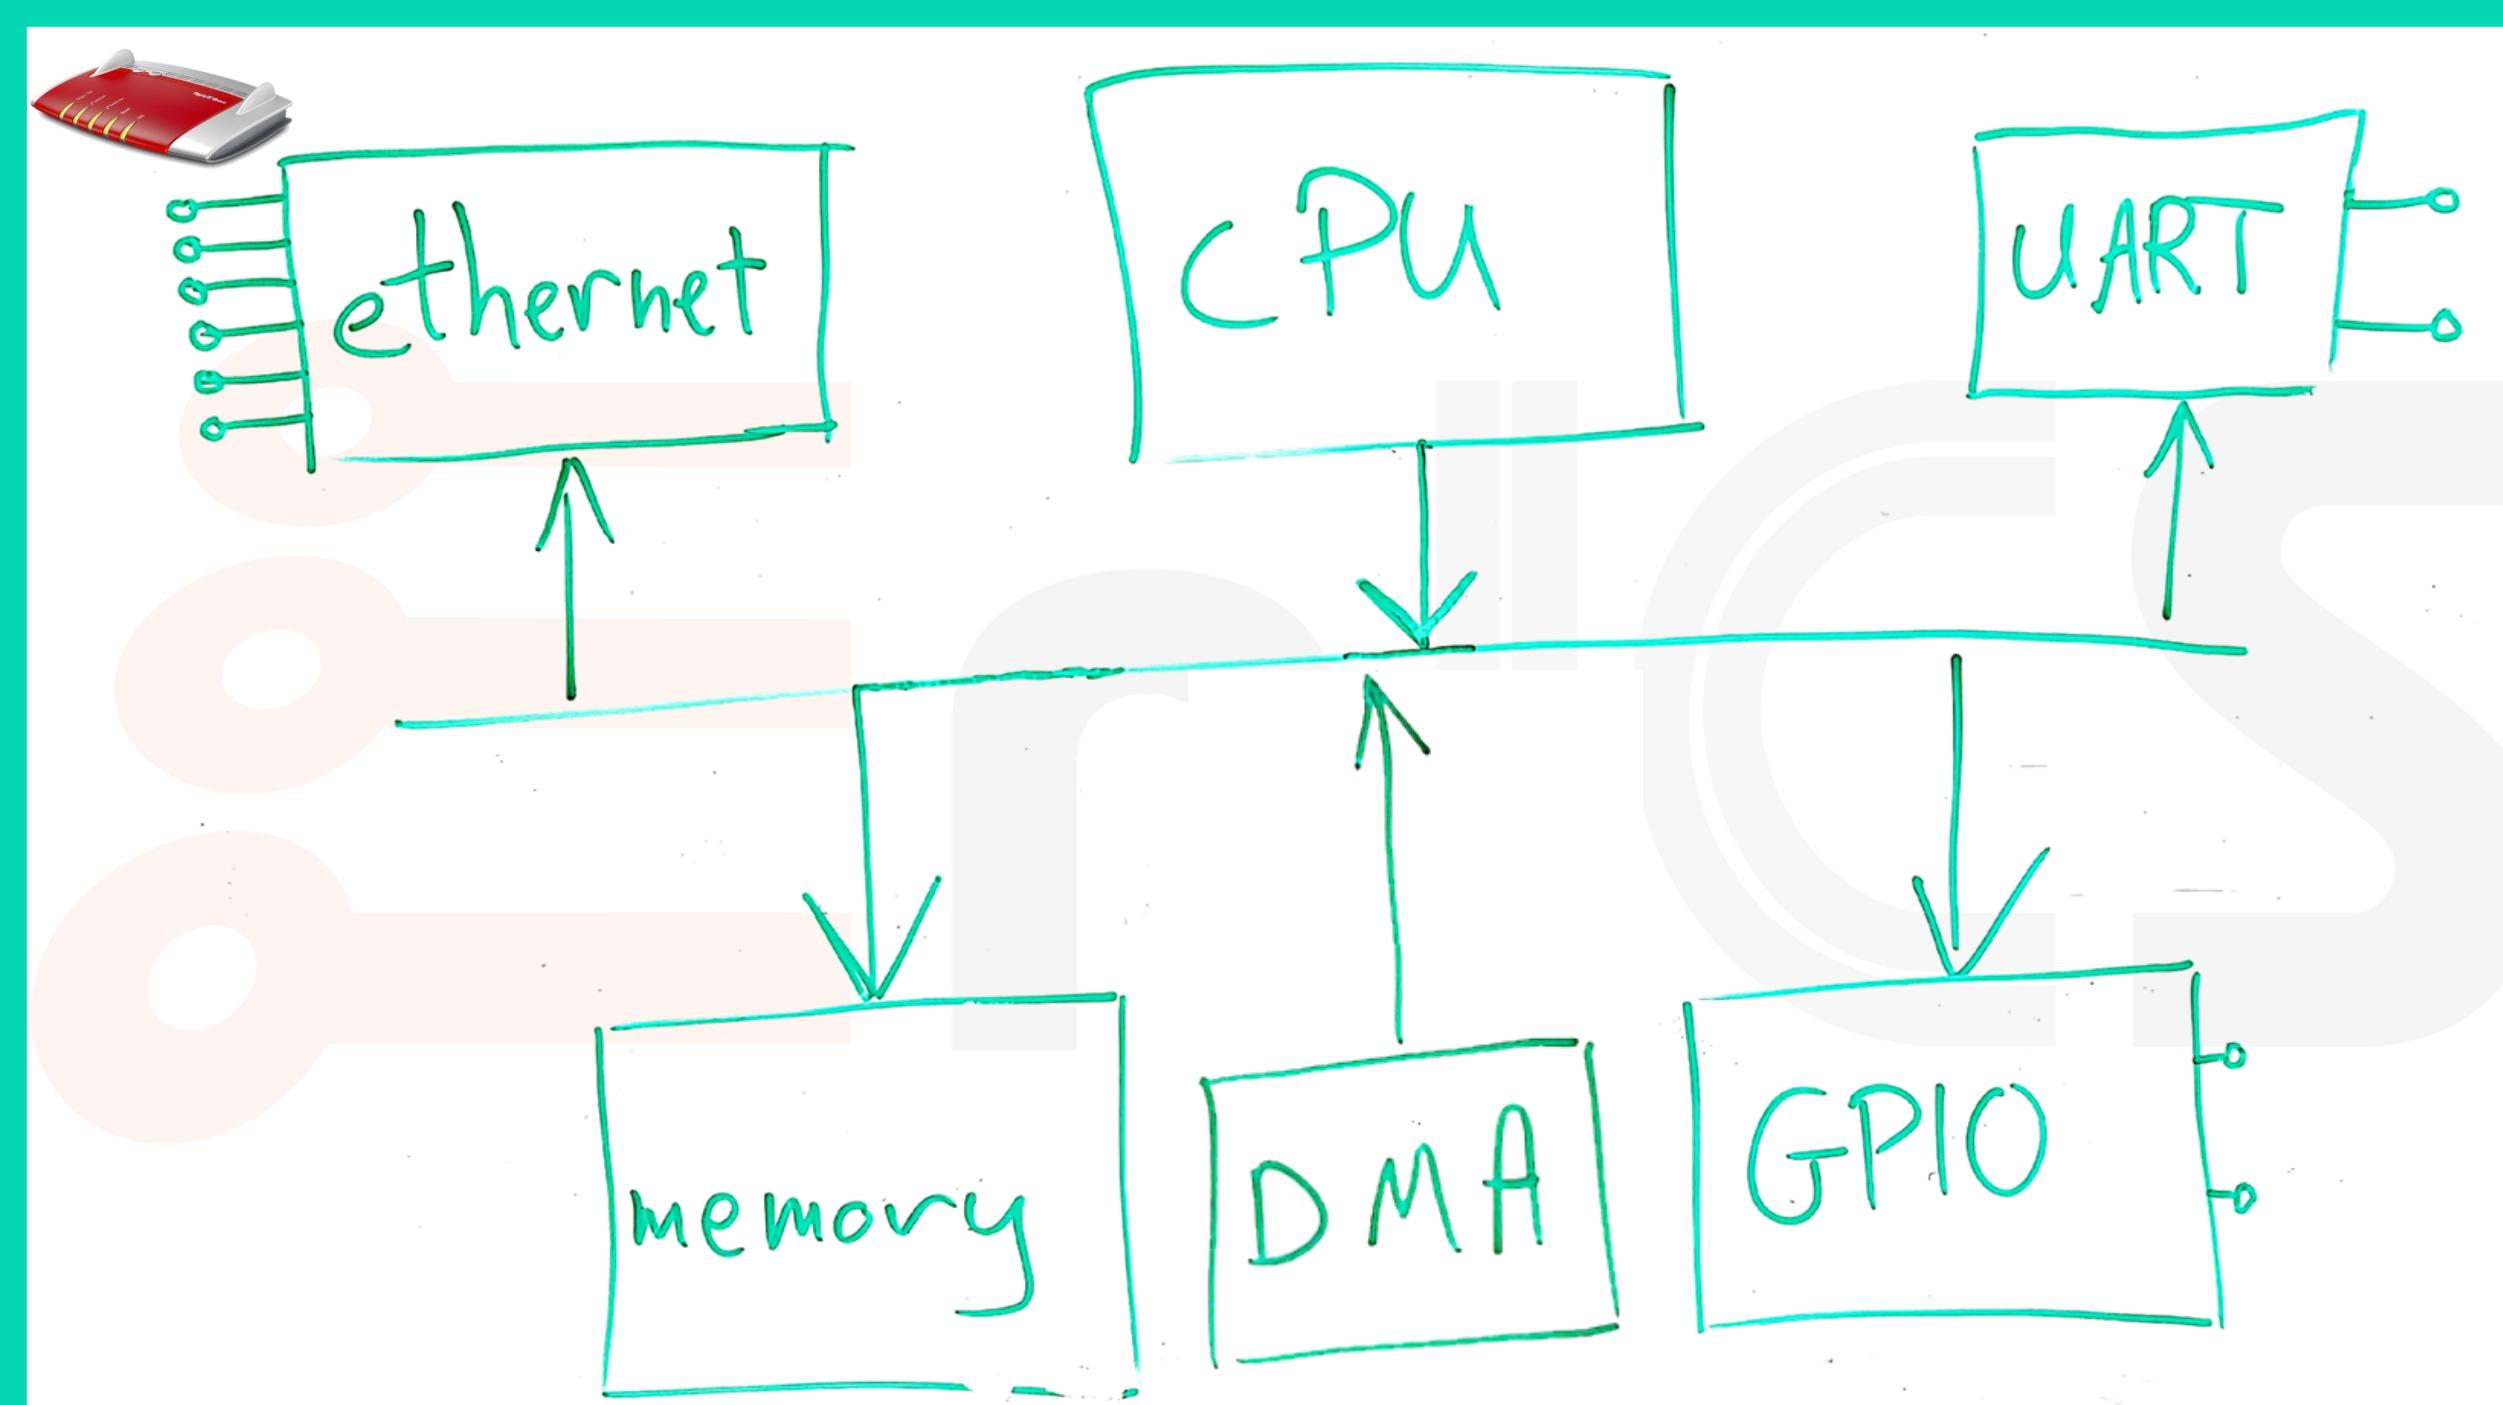
$10^8$  cycles/sec

**CPU THROUGHPUT:**

$1.1 \times 10^9$  bits/sec



So the CPU is again too slow...



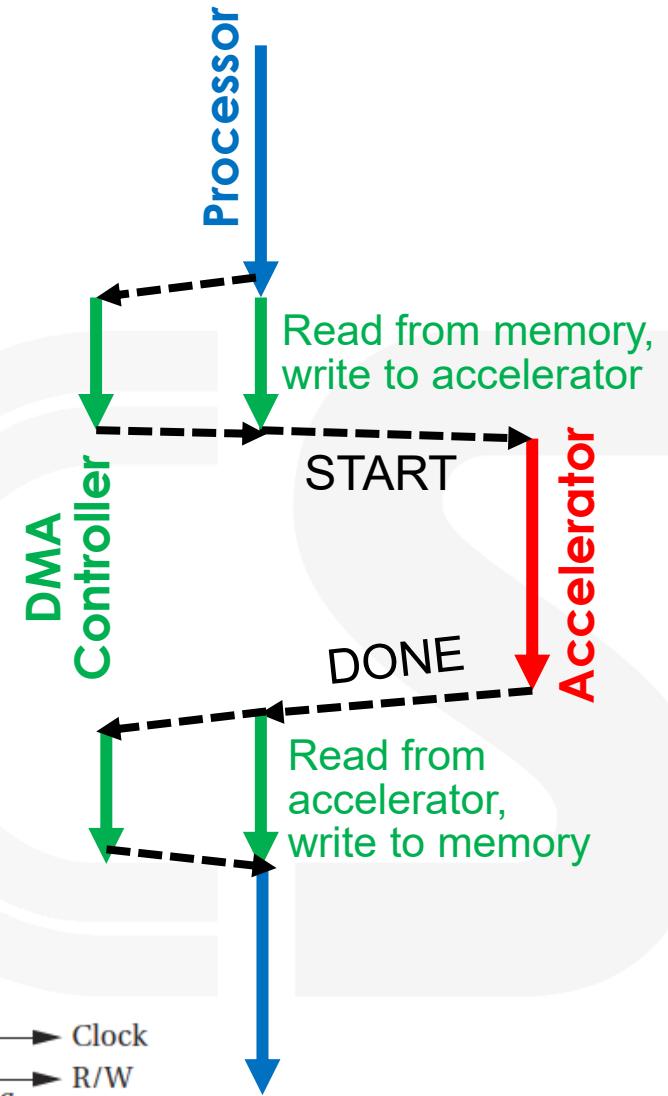
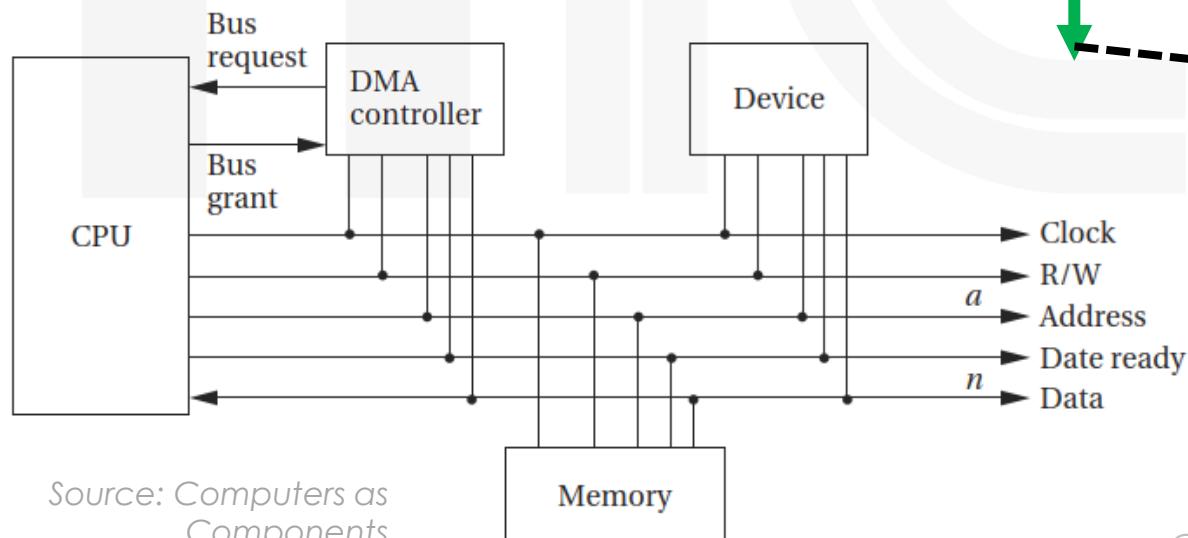
# DMA

- Direct memory access (DMA) is a bus operation that allows reads and writes not controlled by the CPU.

- A DMA transfer is controlled by a DMA controller that requests control of the bus from the CPU.
- With control, the DMA controller performs read and write operations directly between devices and memory.

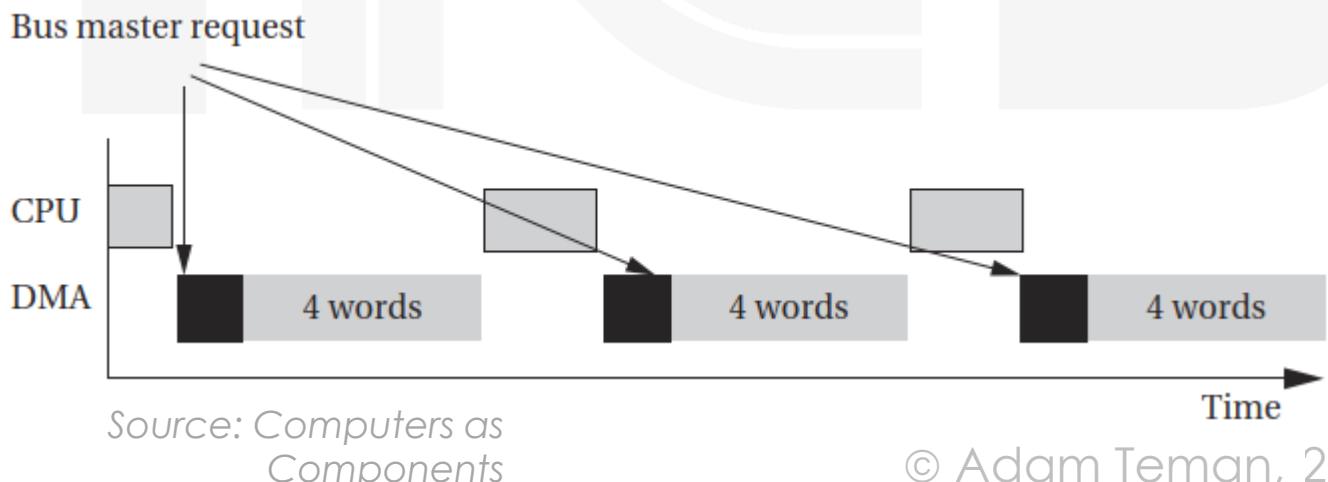
- DMA adds new signals:

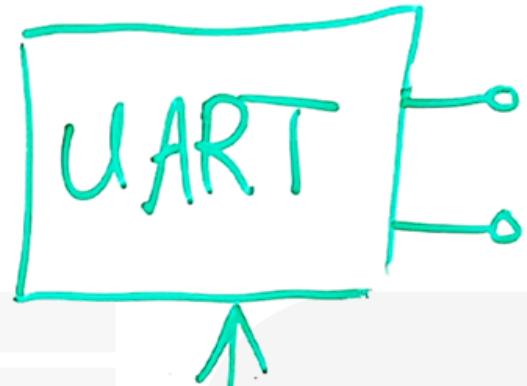
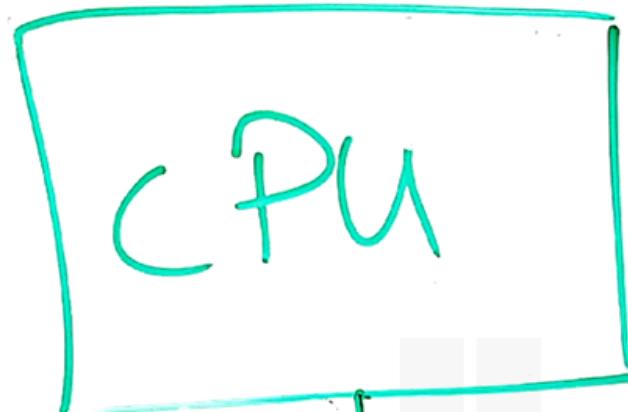
- Bus request
- Bus grant



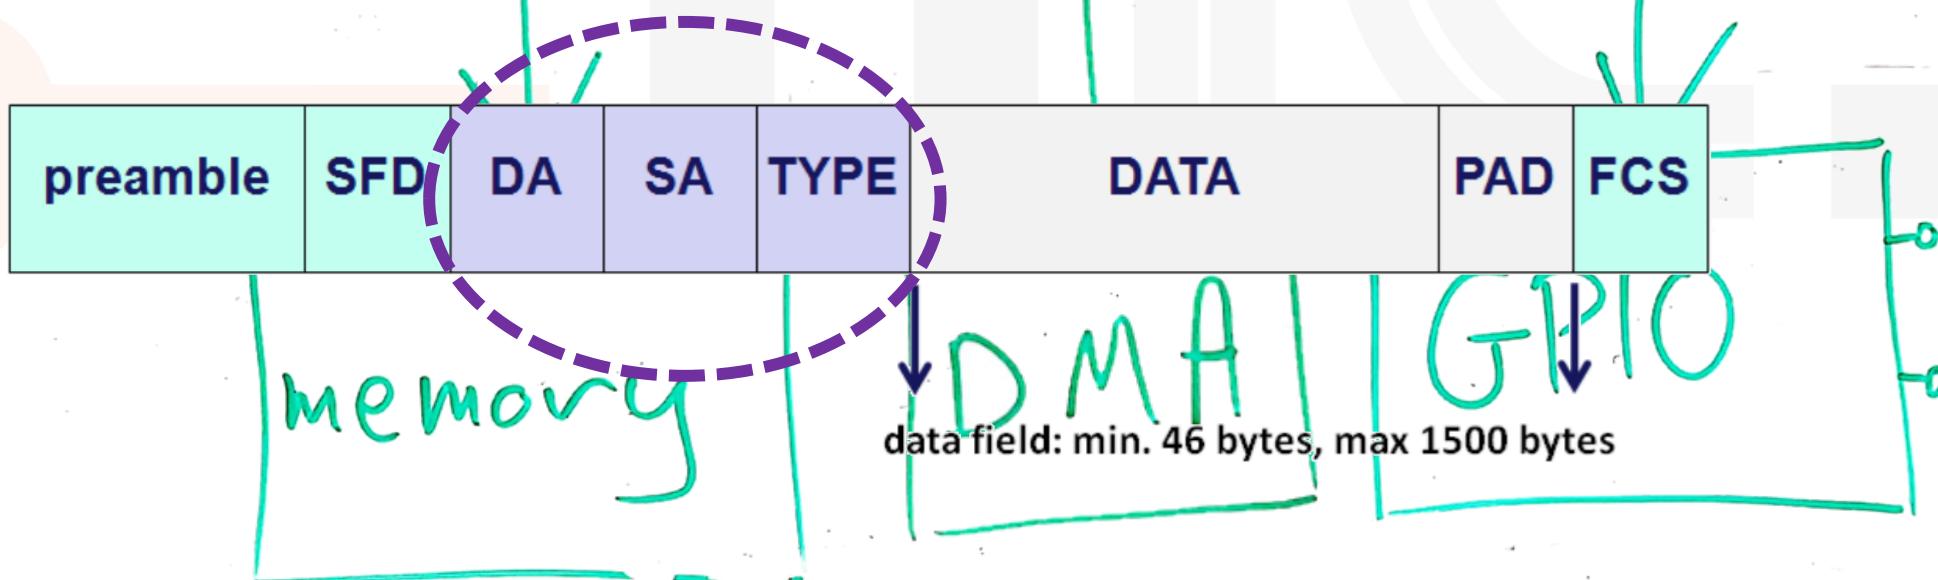
# DMA Registers

- The CPU controls the DMA operation through registers in the DMA controller.
  - Starting address register
  - Length register
  - Status register – to start and stall the DMA
- After the DMA operation is complete, the DMA controller **interrupts** the CPU to tell it that the transfer is done.
- DMA controllers usually use short **bursts** (e.g., **4-16 words**) to only occupy the bus for a few cycles at a time





# BUT WHAT ABOUT THE CPU?



# Data Rate or Packet Rate?

**BIT RATE:**

$2 \times 10^9$  bits/sec

each packet includes 20 bytes of overhead



**DATA RATE:**

Larger packets mean  
interconnect is busier

$\text{BIT-RATE} \times P / (P+20)$

$P=64 \rightarrow 2 \times 0.76 \times 10^9$  bits/sec

$P=1518 \rightarrow 2 \times 0.98 \times 10^9$  bits/sec

**PACKET RATE:**

Smaller packets mean  
CPU has to do more

$\text{BIT-RATE} / ((P+20) \times 8)$

$P=64 \rightarrow 2 \times 1.48 \times 10^6$  packets/sec

$P=1518 \rightarrow 2 \times 81.2 \times 10^3$  packets/sec

# How does this affect the CPU?

- The CPU (in a router) needs to handle the packet
  - i.e., figure out where to send the packet to.
  - So, all it cares about is *packet rate*
- How much work can the CPU do on each packet?
  - For packets with 1518 bytes of data, the packet rate is about 160K packets/sec
  - At 100 MHz → 615 instructions per packet → Not that much
  - For 64 byte packets → 34 instructions per packet → Infeasible!
- What can we do???
  - Trivial solution: Raise the frequency
  - Still not enough: Add additional CPUs
  - Better solution: Integrate dedicated hardware (Accelerators and ASIPs)



SPEED

# A typical router SoC



GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

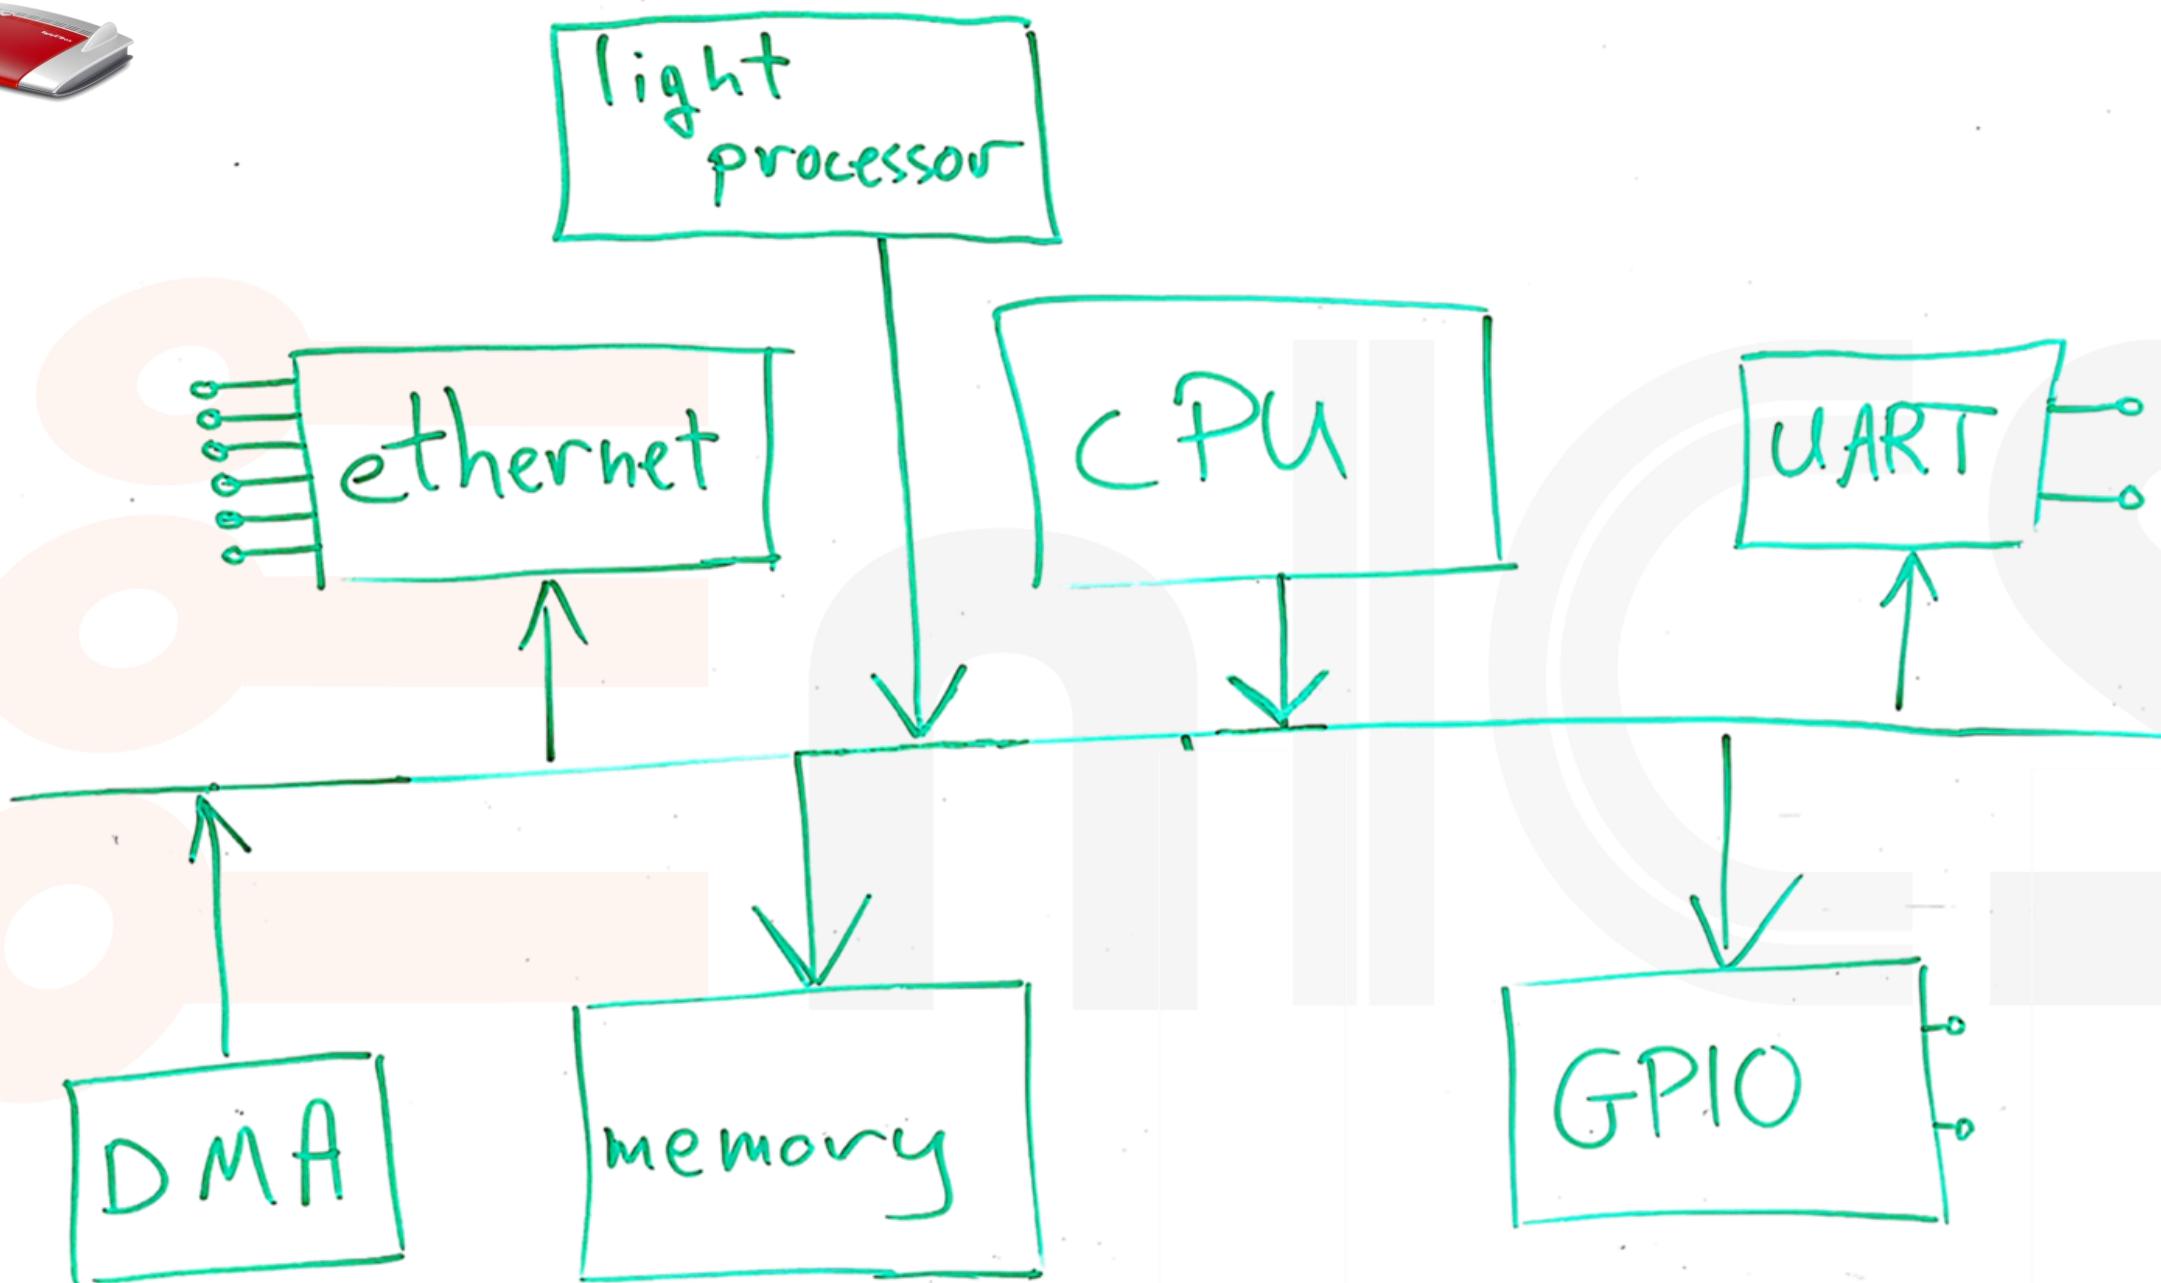
# Memory



Emerging Nanoscaled  
Integrated Circuits and Systems Labs

The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University





# But what about memory?

- We now have a processor that can communicate with peripherals, with an off-chip network, etc.

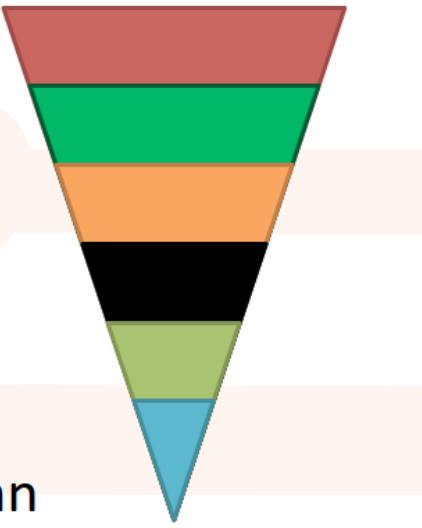
*But what about memory?*

- Our router needs a lot of memory:
  - To **buffer** packets
  - To **store** routing tables
  - To **host** the operating system
  - ...
- The on-chip memory (**~MB**) is nowhere near enough.

**We need to use DRAM**

# DRAM Organization

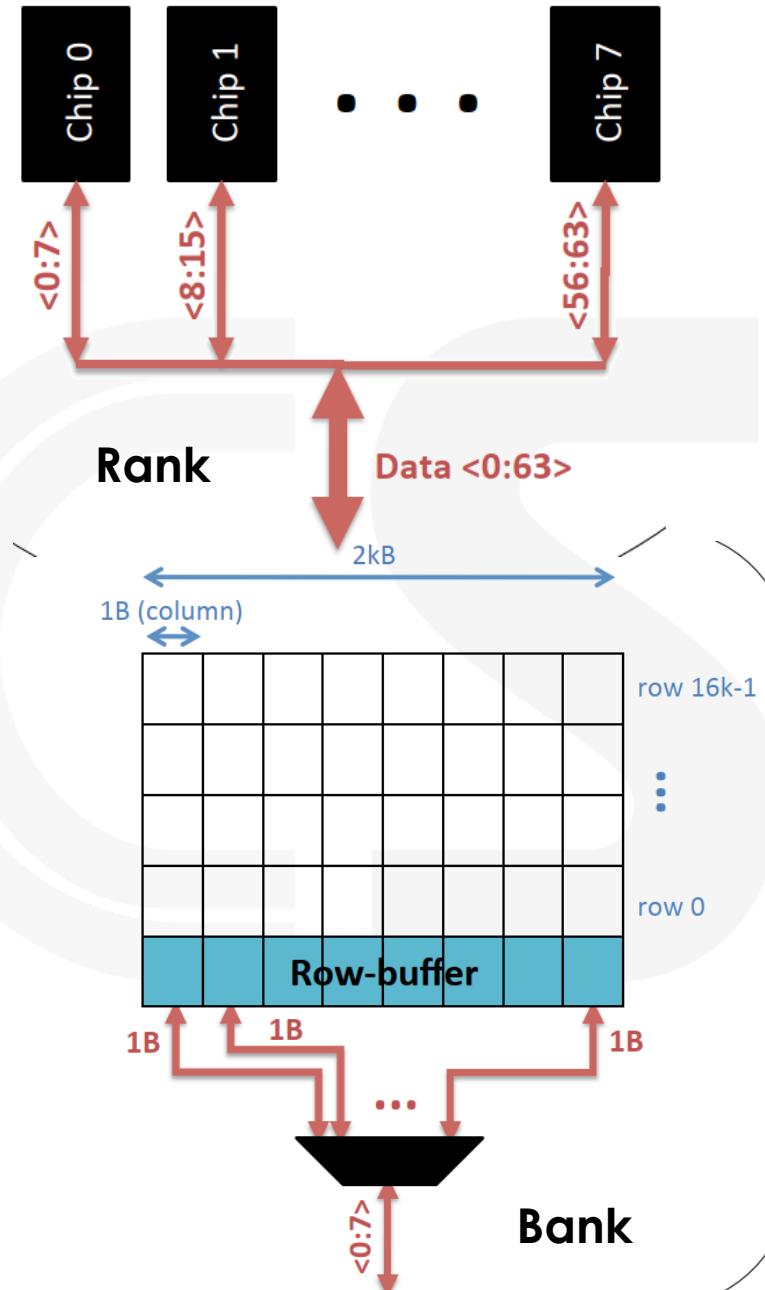
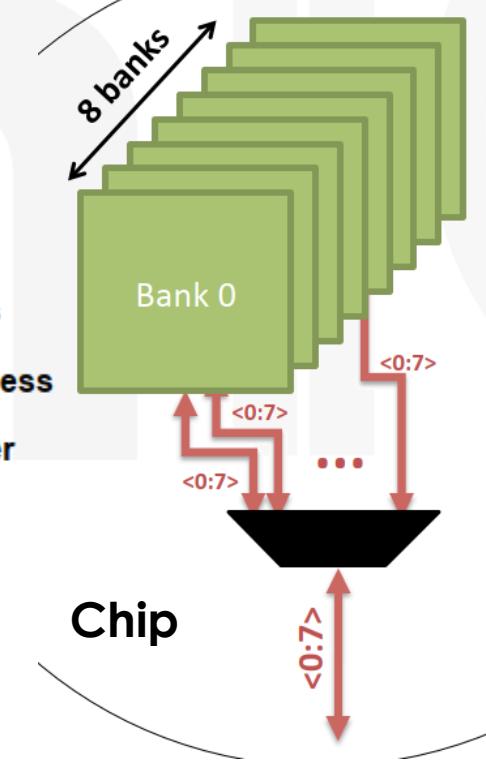
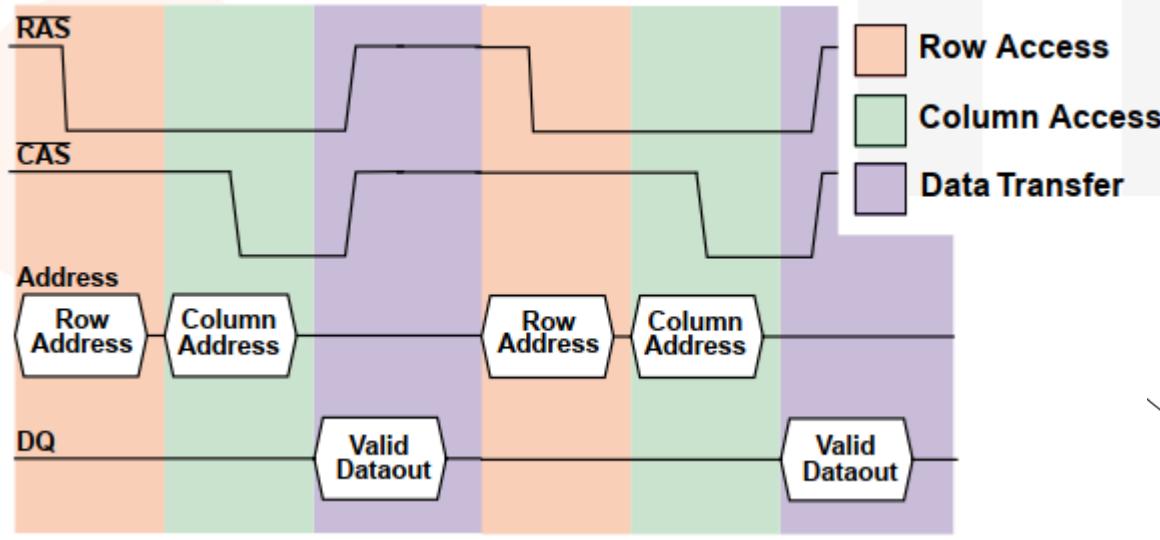
- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column



DIMM (Dual in-line memory module)



Source: Onur Mutlu



# DRAM Organization

- So, for a **1GB DIMM** (i.e., 8 chips),  
we need chips with **1Gb** of memory

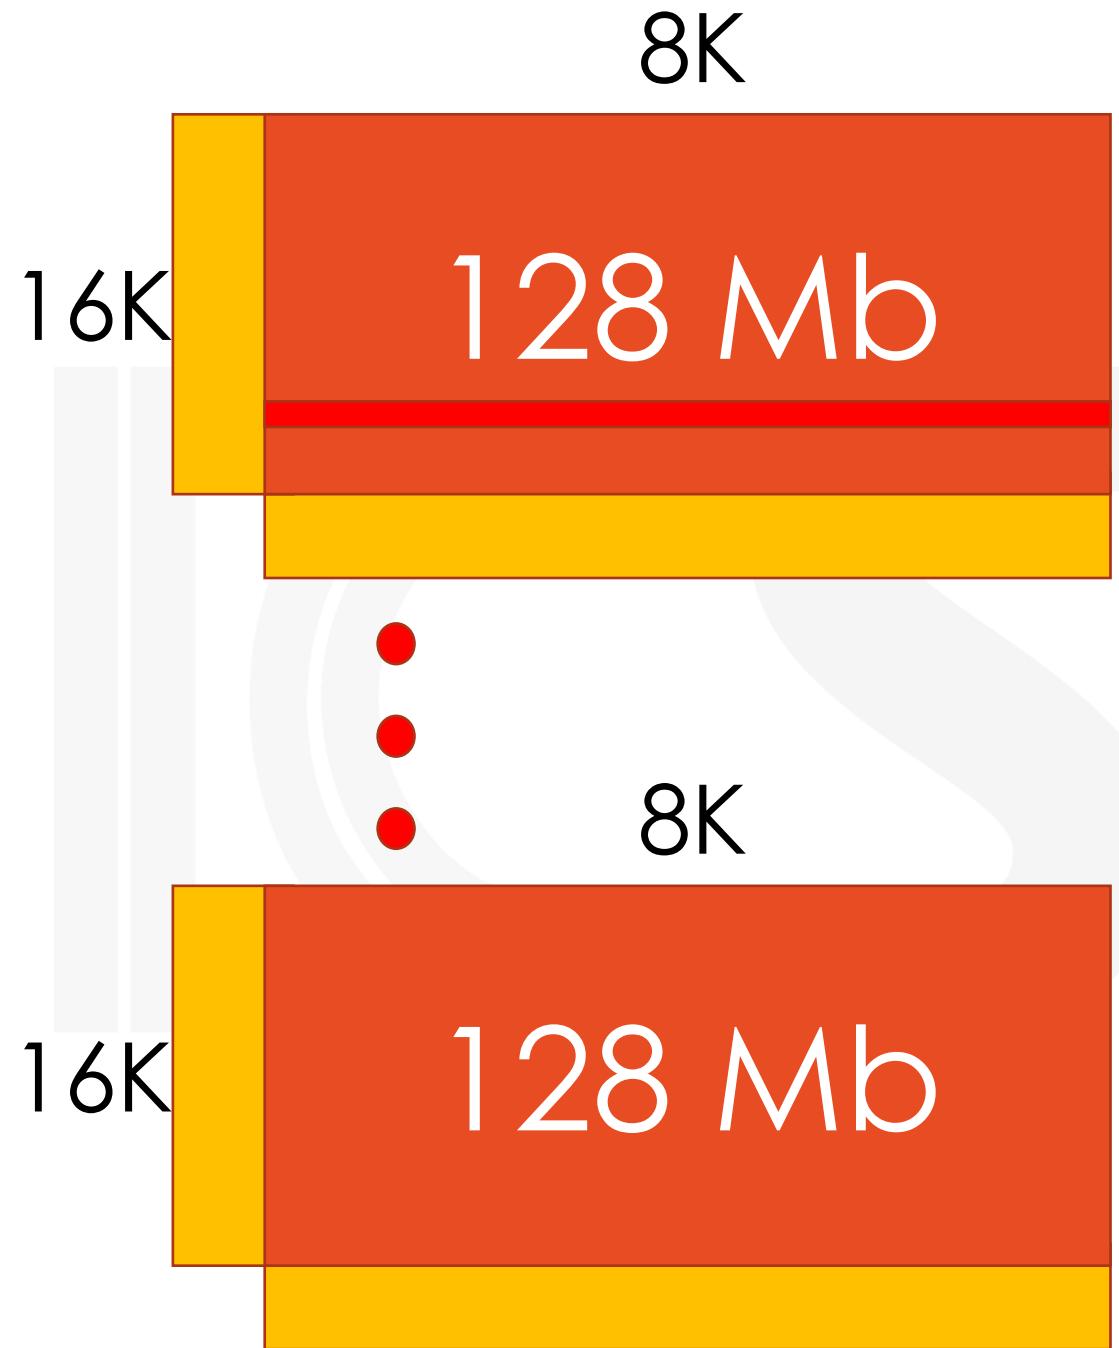
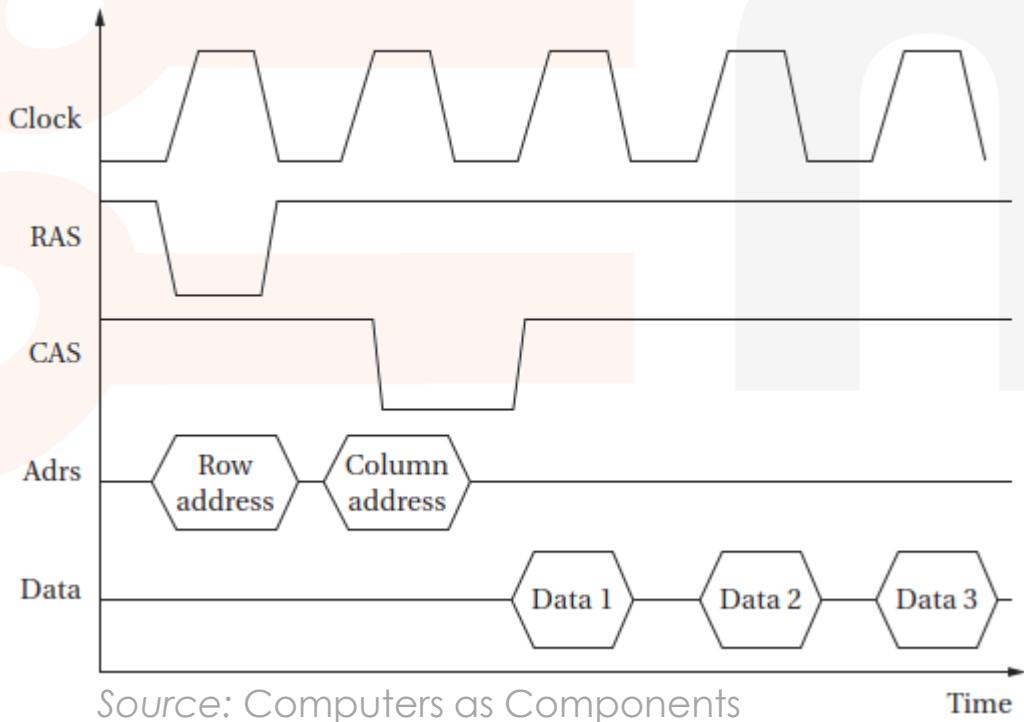
- e.g., **128k x 8k**

- But that is a lot of rows...



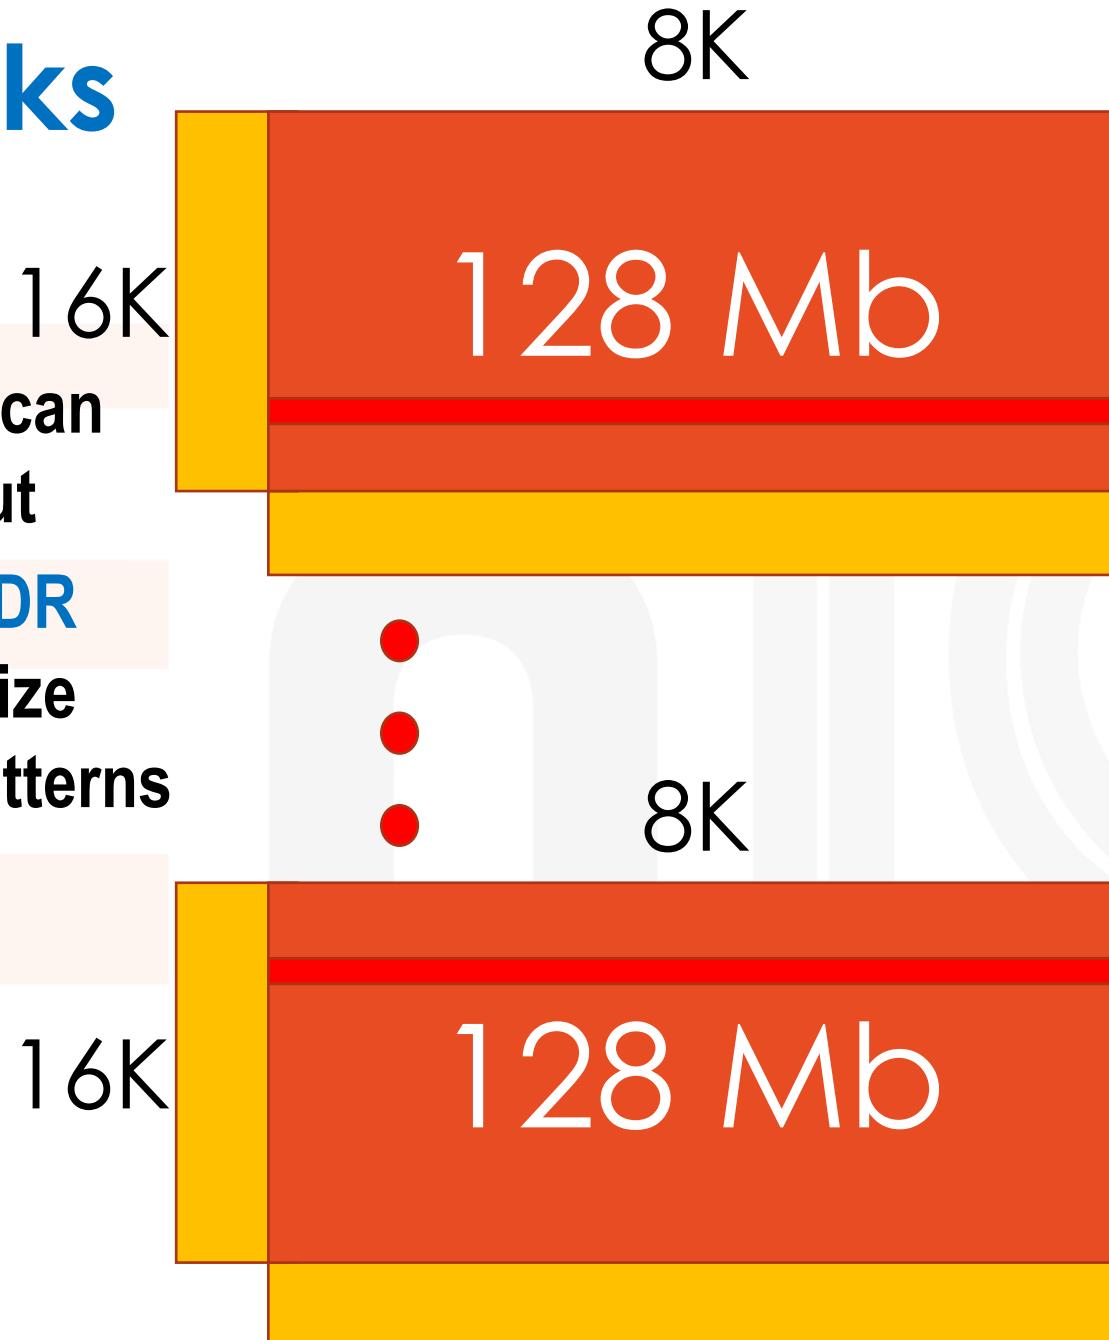
# DRAM Banks

- So, we break it into 8 banks of 16k rows
  - And readout an entire row to a buffer
  - Once the row is buffered, we can directly access any byte in the buffer.



# DRAM Banks

- Reordering the memory accesses can improve throughput
- Add a dedicated **DDR controller** to optimize memory access patterns



BO/R0  
B1/R1  
BO/R0  
B1/R1  
BO/R0  
B1/R1  
BO/R3  
BO/R0  
B1/R2  
B1/R1

BO/R0  
B1/R1  
BO/R3  
B0/R0  
B1/R1  
B0/R3  
B1/R2  
B1/R1

GPIO and UART

Accelerators

Ethernet

DMA and ASIP

DDR

AXI and Boot

# Finishing our Design



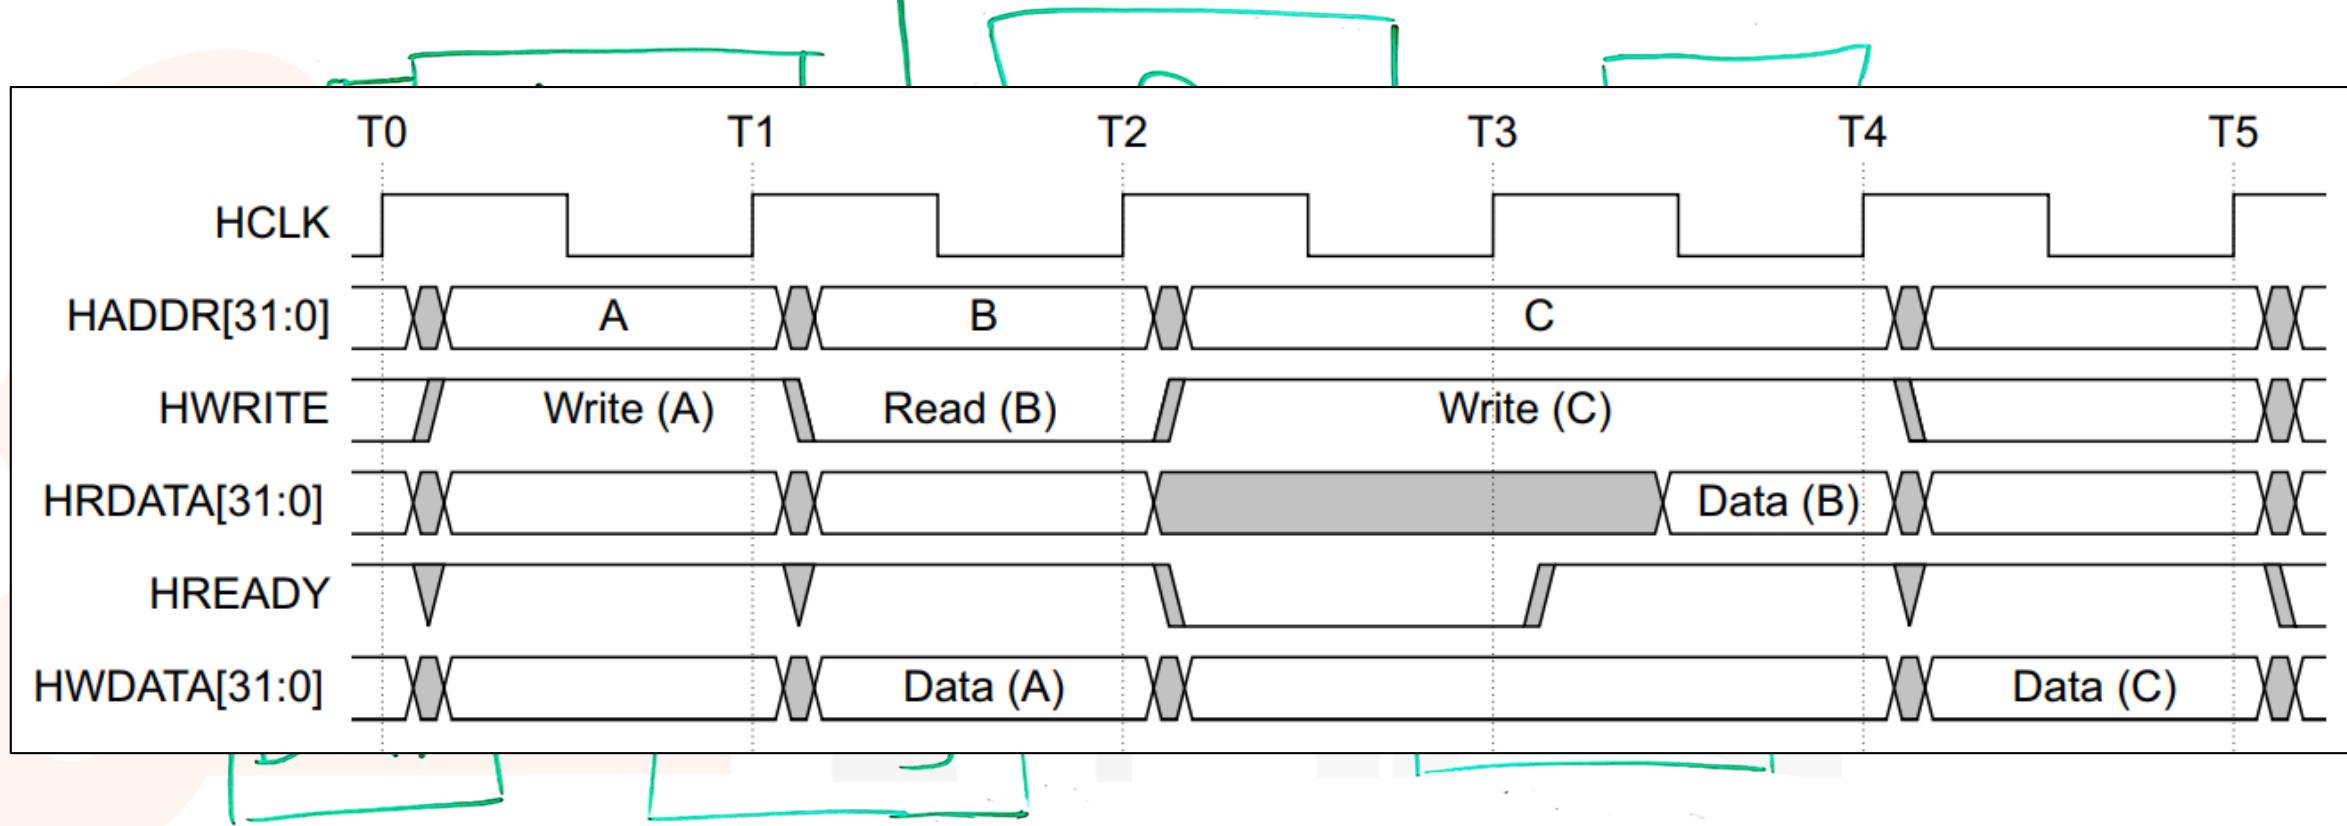
Emerging Nanoscaled  
Integrated Circuits and Systems Labs

The Alexander Kofkin  
**Faculty of Engineering**  
Bar-Ilan University





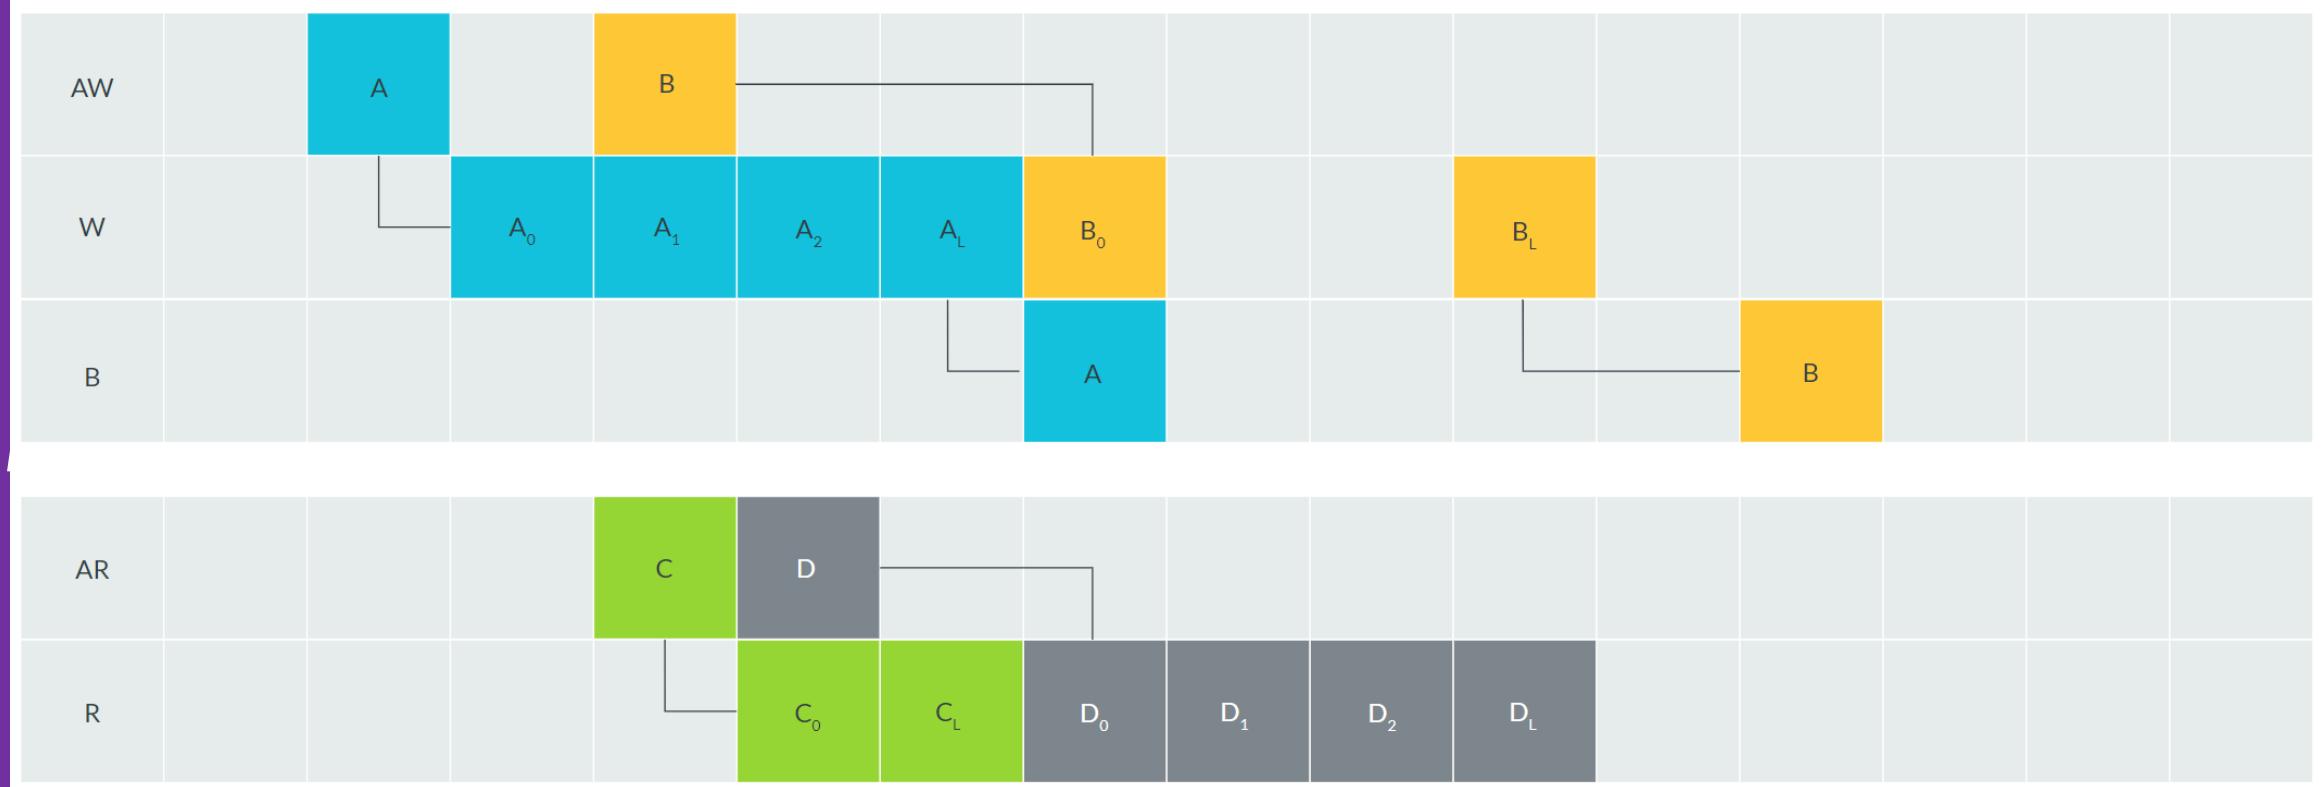
light  
processor



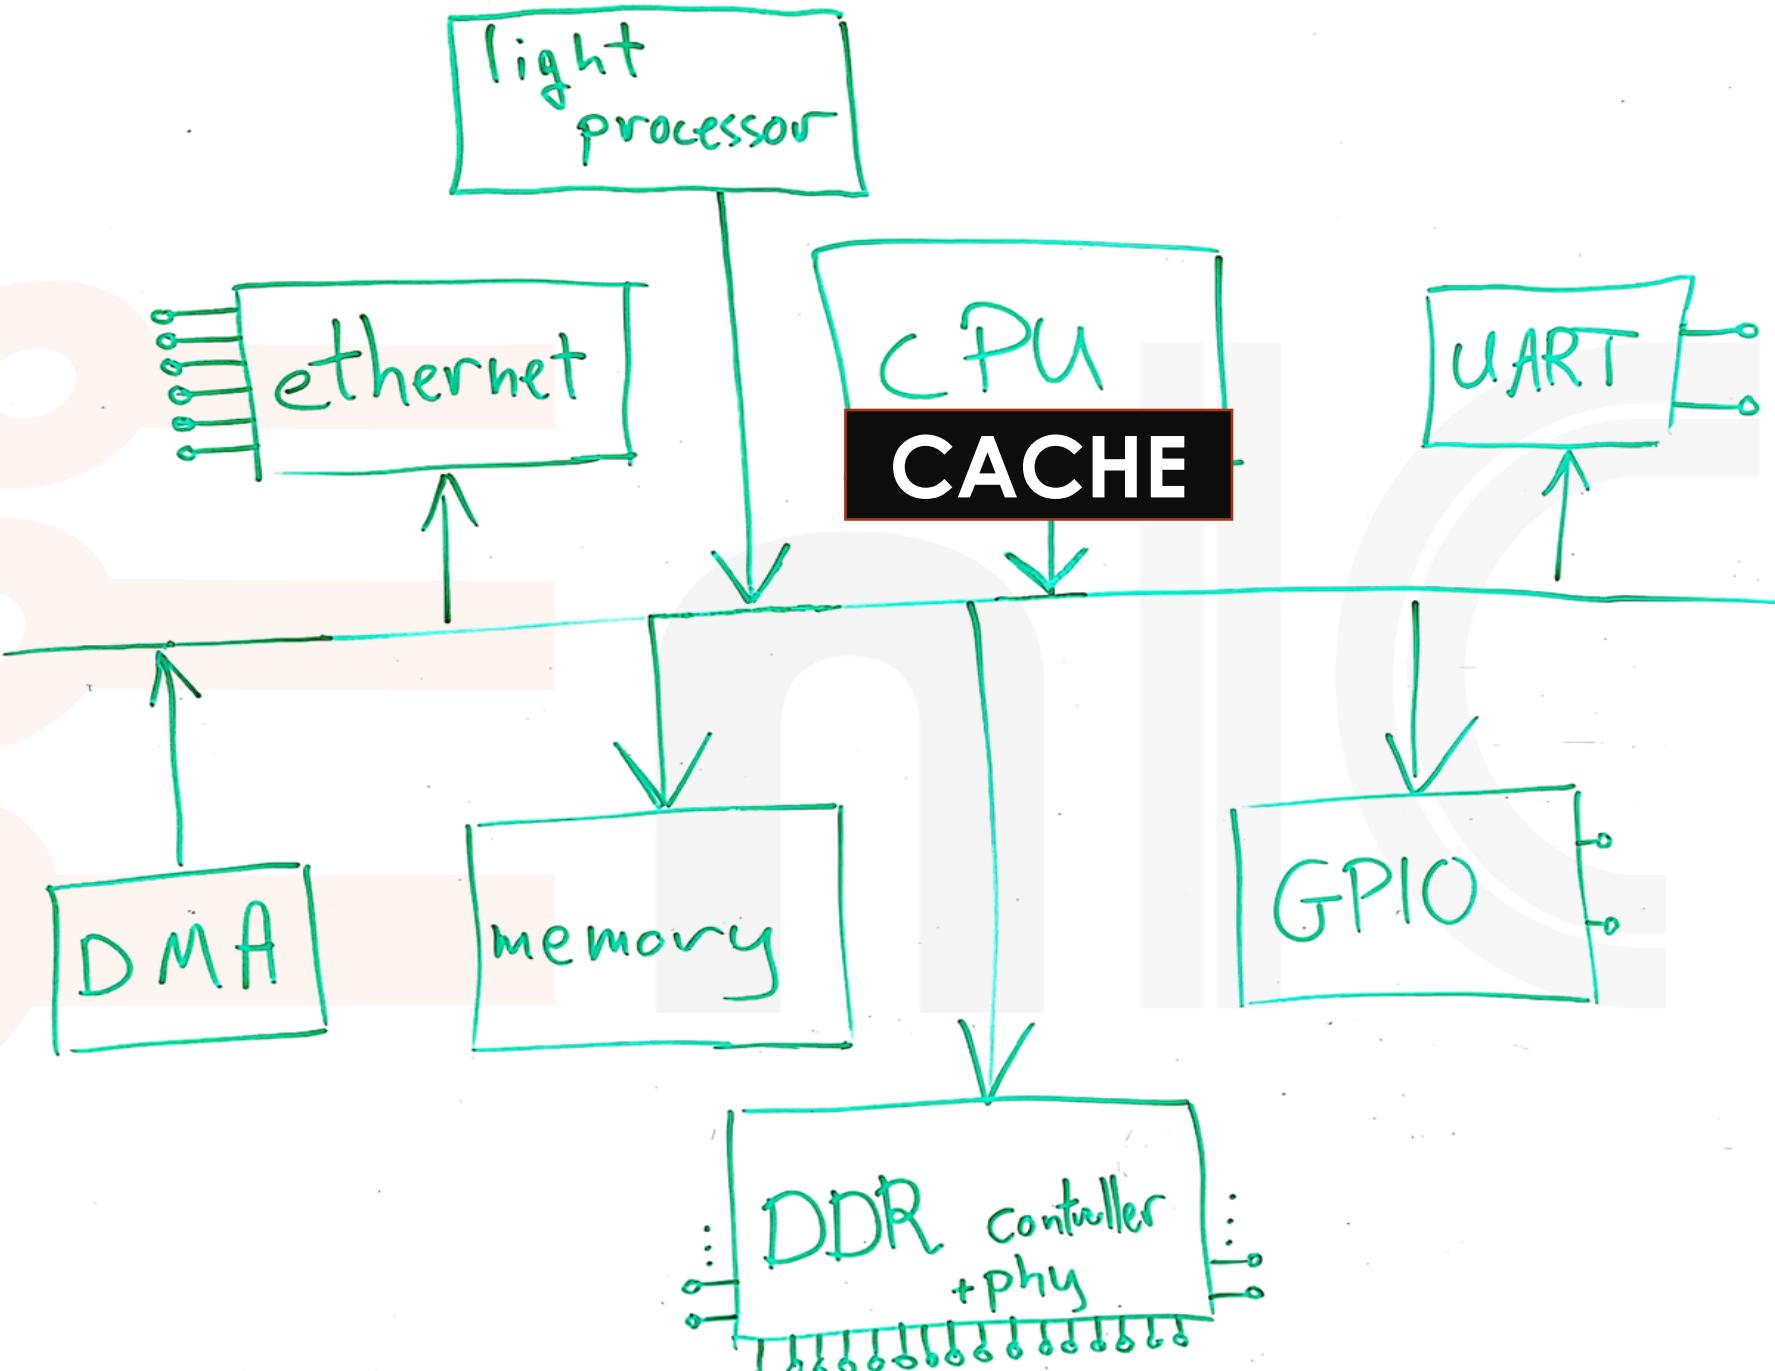
In AHB, we always have to wait until a transaction is finished before starting a new one...

# AXI

WRITE ADDR CHANNEL



READ DATA CHANNEL

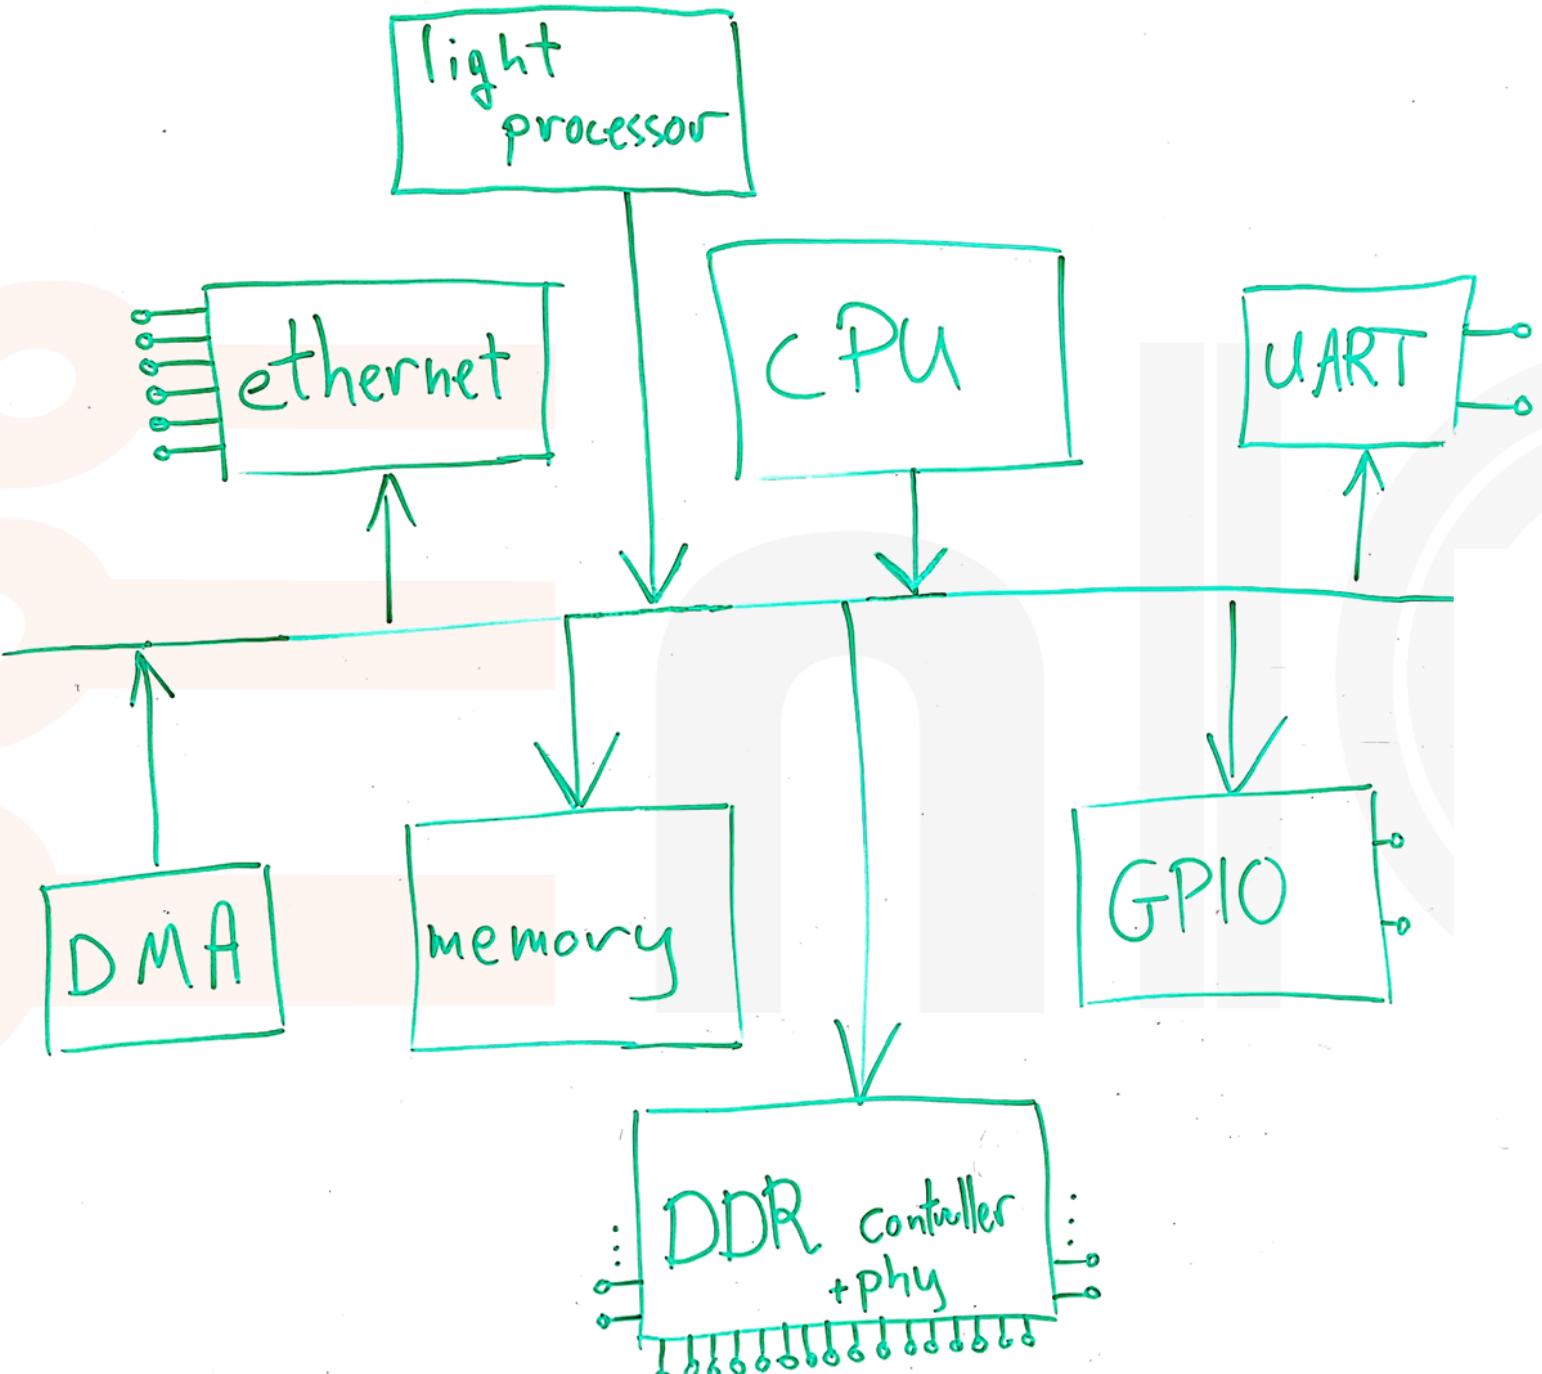


# But what happens on Startup?

- Start by reading from a **BootROM**
  - A small piece of memory, which contains the very first code that is executed upon reset.
  - Either hard-wired (**mask-ROM**) or rewriteable (**embedded Flash**)
  - Can use **bootstraps/fuses** to change configuration.
- Then move on to the **Bootloader**
  - Usually stored on rewriteable flash (i.e., **SD card**)
  - Configures the chip and some of the peripherals
  - Loads the end application (e.g., **OS**) from storage (**flash, SSD, HDD**)
  - Passes control to the end application
- The **BootROM** and **bootloader** can be combined
- The **BootROM** of an **x86 system** is called the **BIOS**

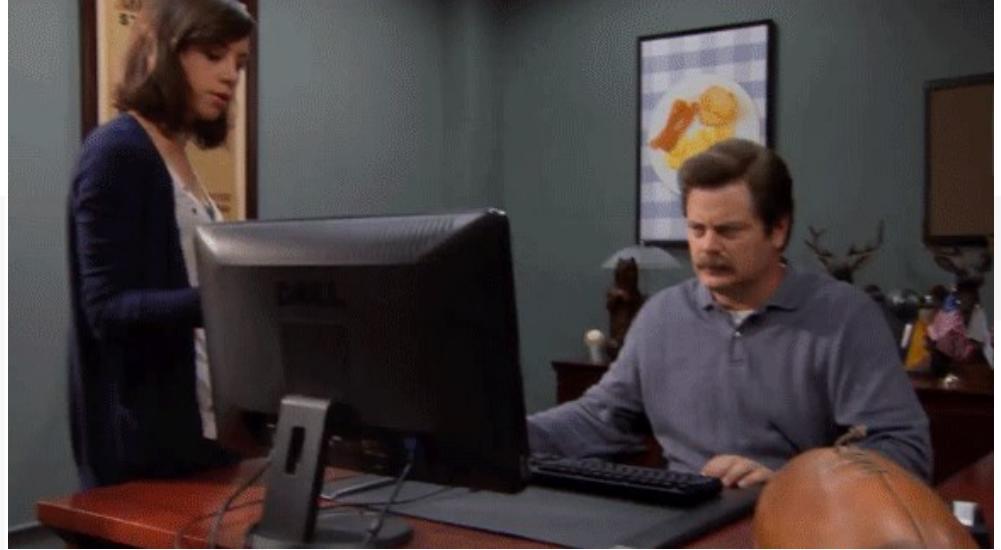


Source: Intel

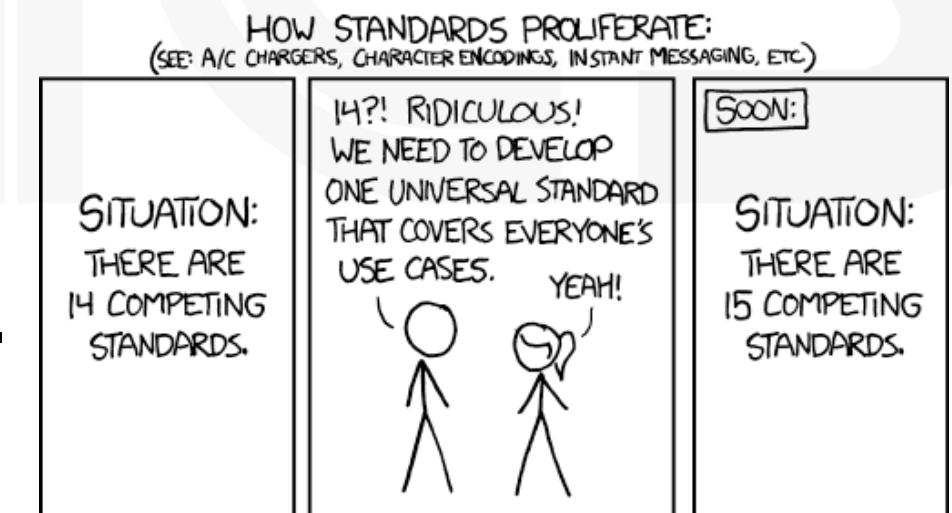


# Conclusions

- Processors are great at processing.
- They are not so great at data movement.
- They are not so great at doing simple tasks.
- For a well-defined task, dedicated hardware will never lose to a processor.
- With Processors we gain flexibility.
- Software development is faster than Hardware.
- Bugs in Software are much cheaper to fix.



Parks and Recreation. Source: Giphy



Source: xkcd

© Adam Teman, 2023

# Main References

- Tzachi Noy, “Interfaces: External/Internal, or why CPUs suck”, 2019
- Wolf, “Computer as Components - Principles of Embedded Computing System Design,” Elsevier 2012